

Procedures

Procedures

A **Procedure** is a unit of code that we can call

Depending on the programming language, it may be called a procedure, function, subroutine, or method

A call is like a jump, except it can return.

The hardware provides machine instructions for this:

`call label`

Push return address on stack; Jump to `label`

`ret`

Pop address from stack; Jump to address

First of all, we have to understand how a stack works...

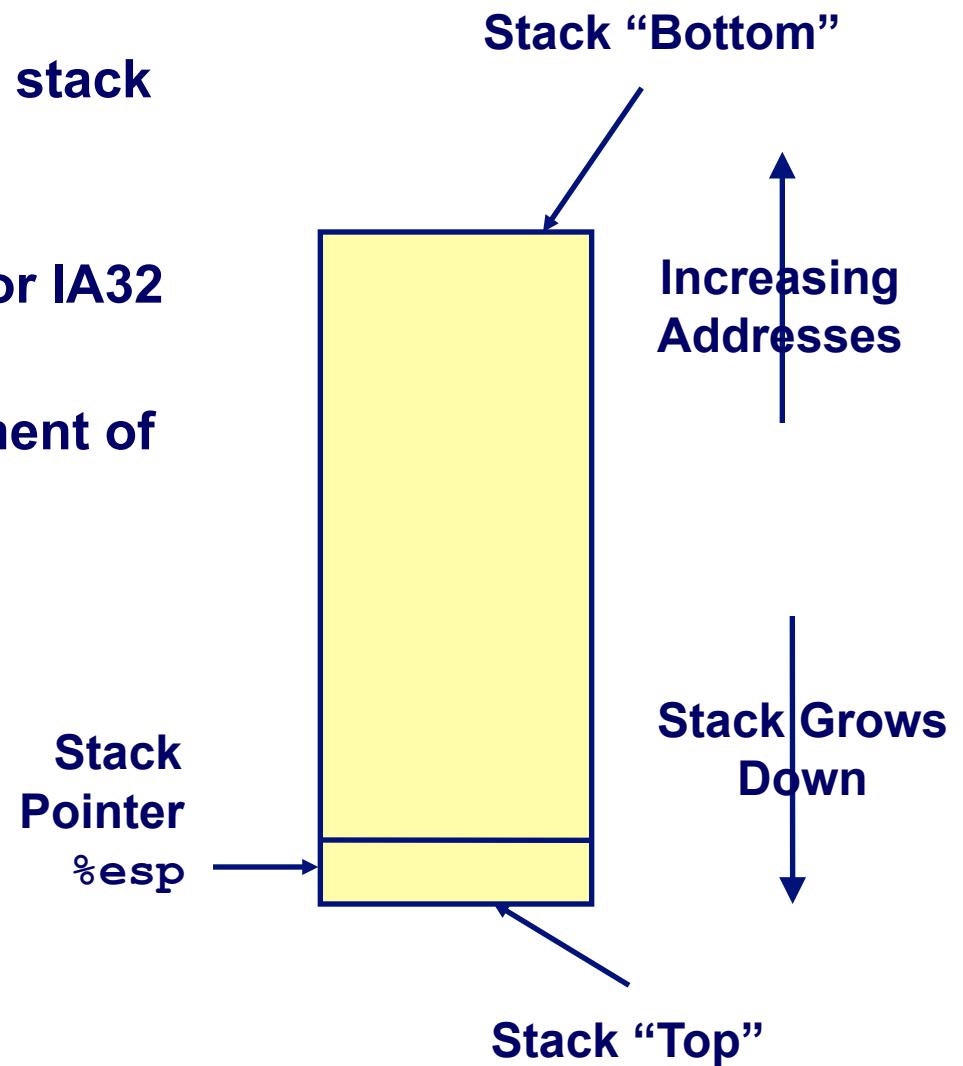
IA32 Stack

Region of memory managed with stack discipline

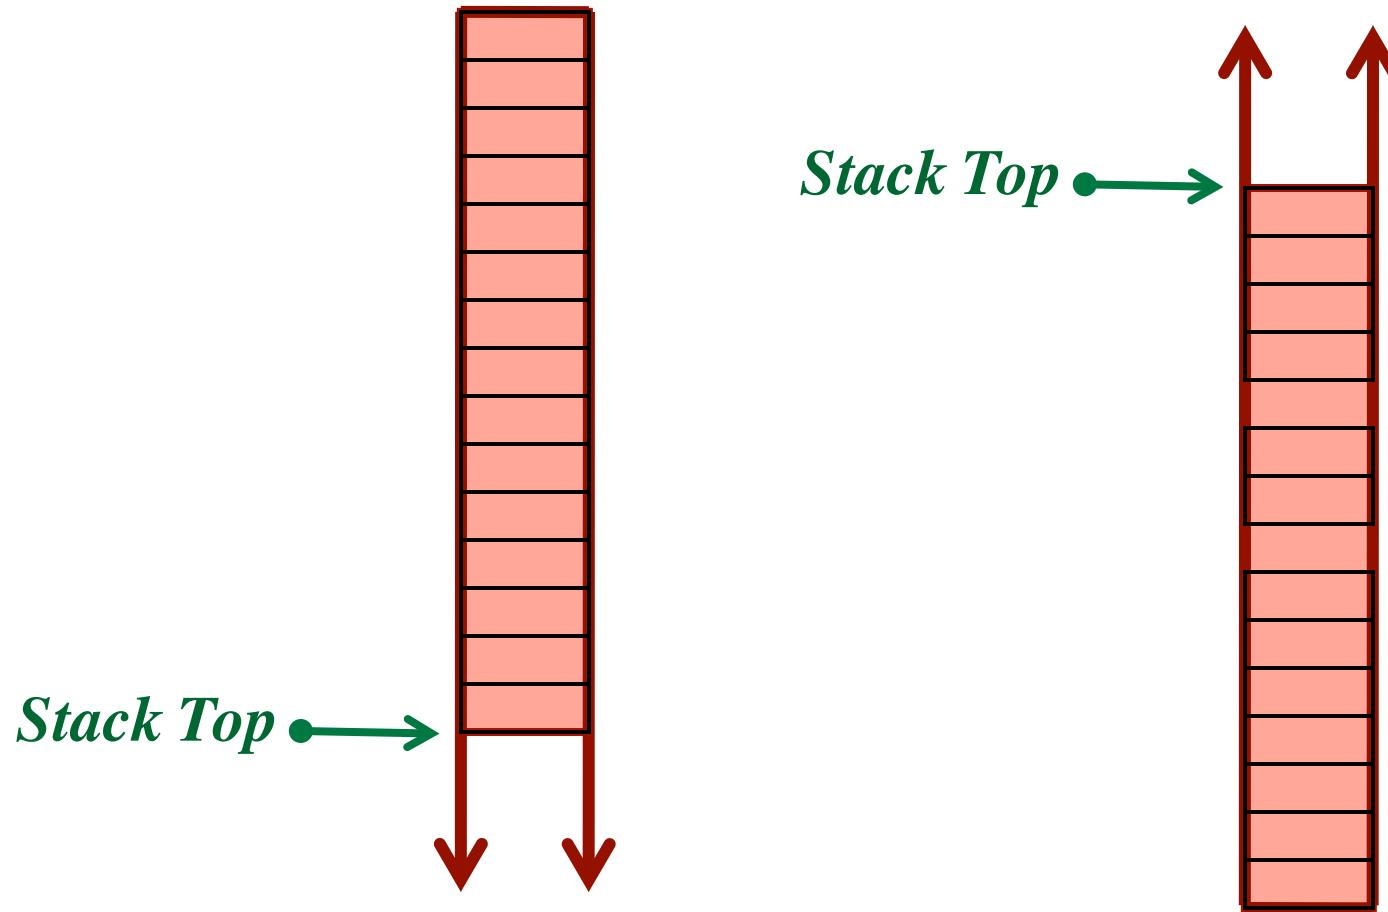
Grows toward lower addresses for IA32

Register `%esp` indicates top element of stack

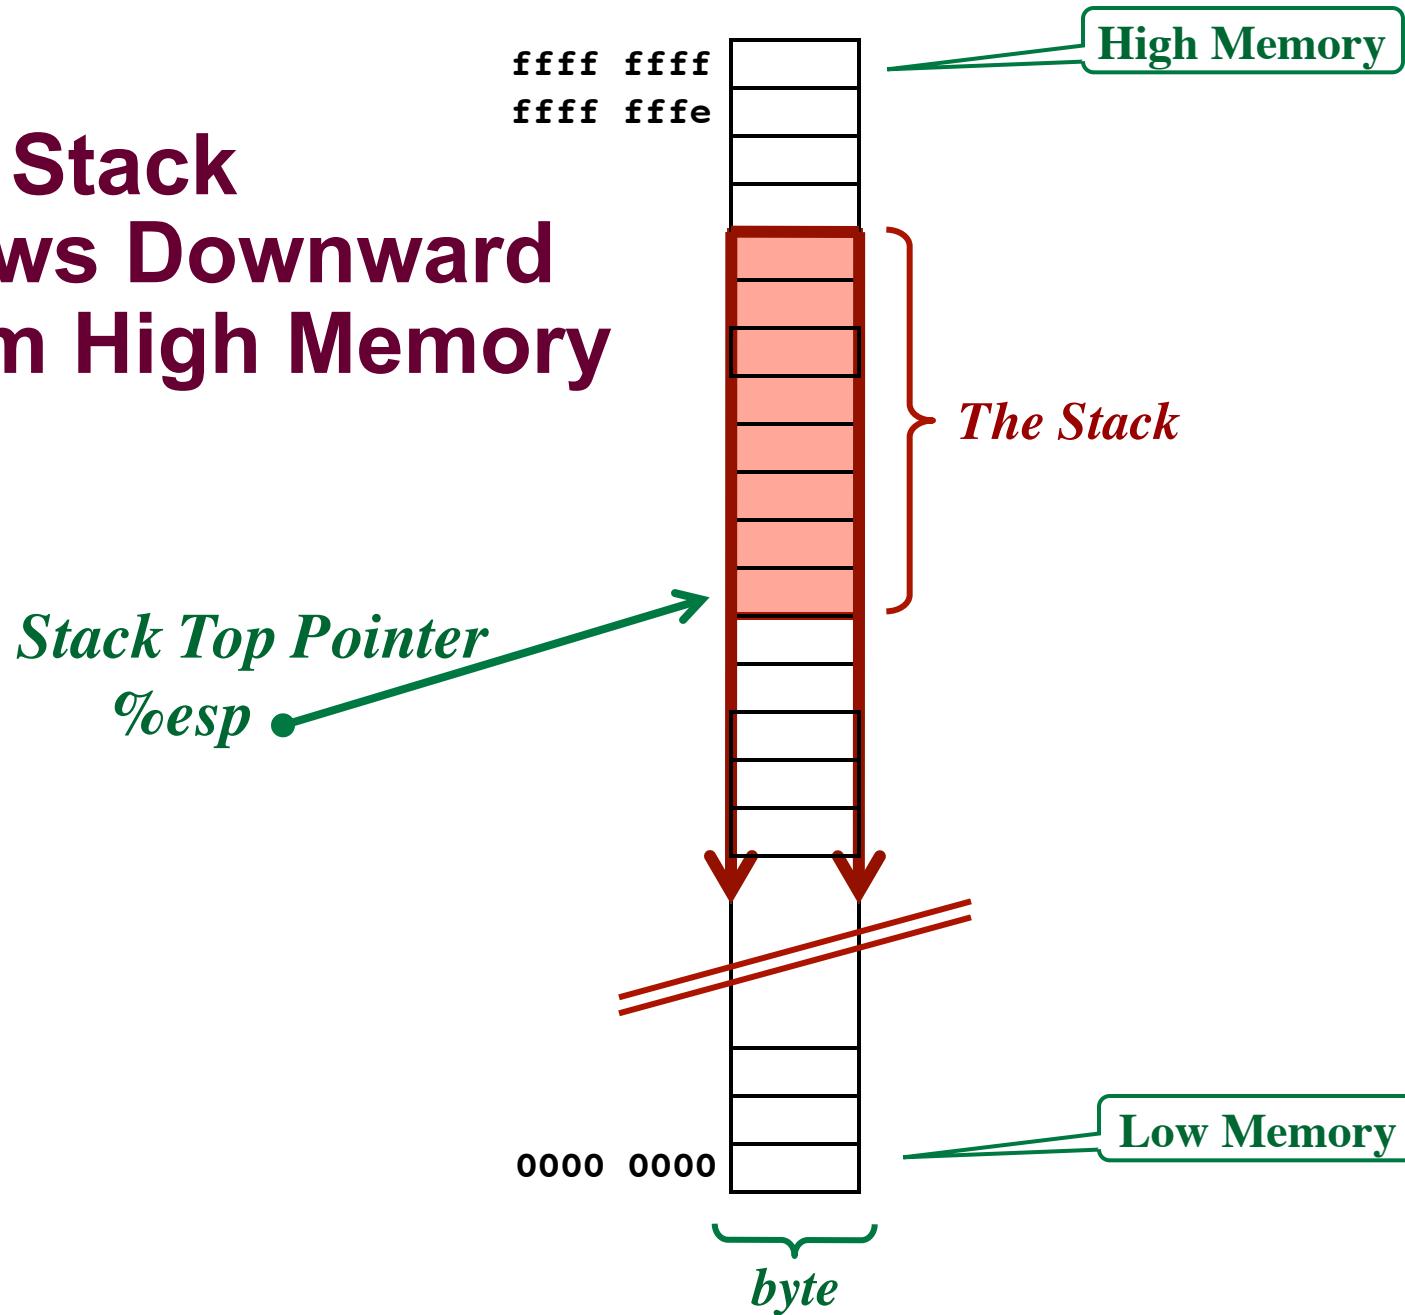
Top element has lowest address



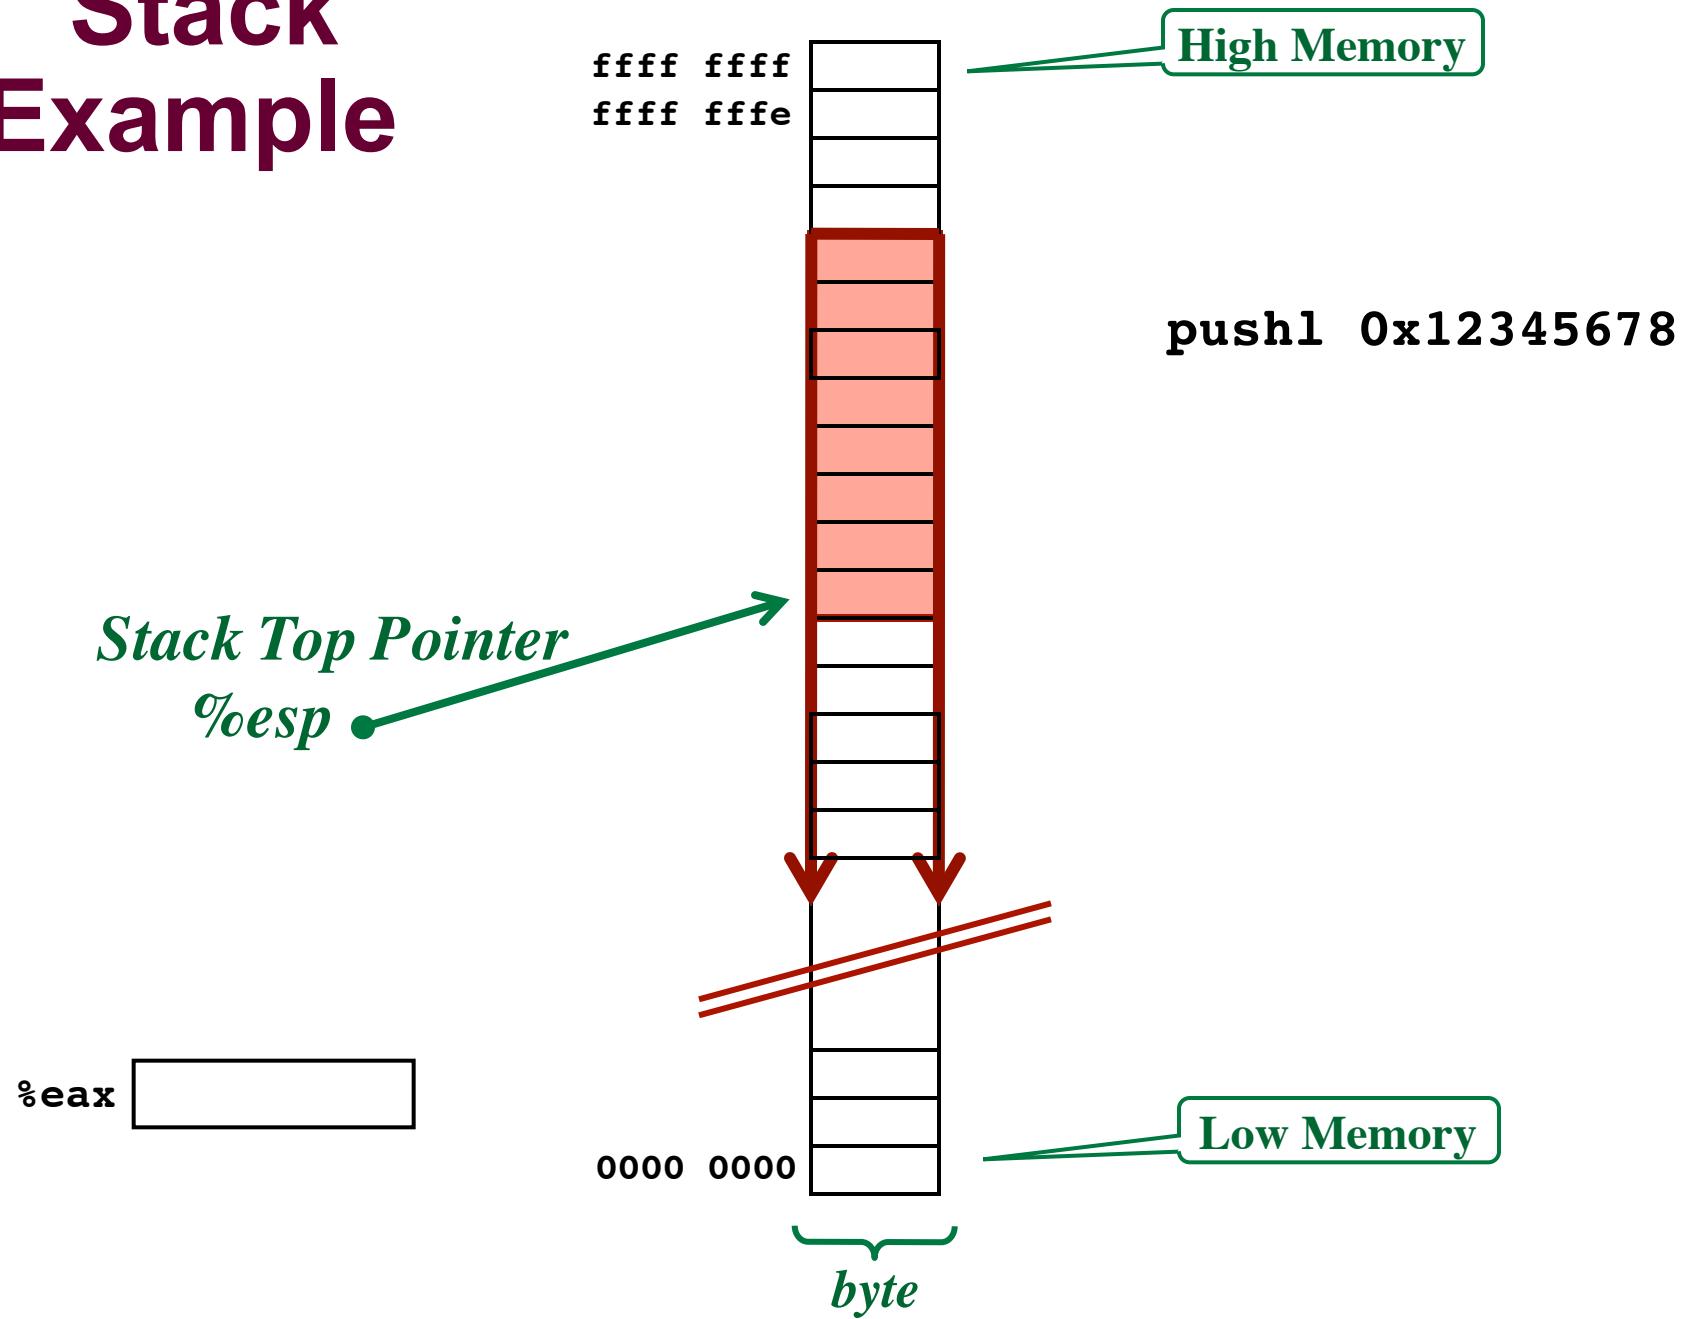
How shall we imagine a stack?



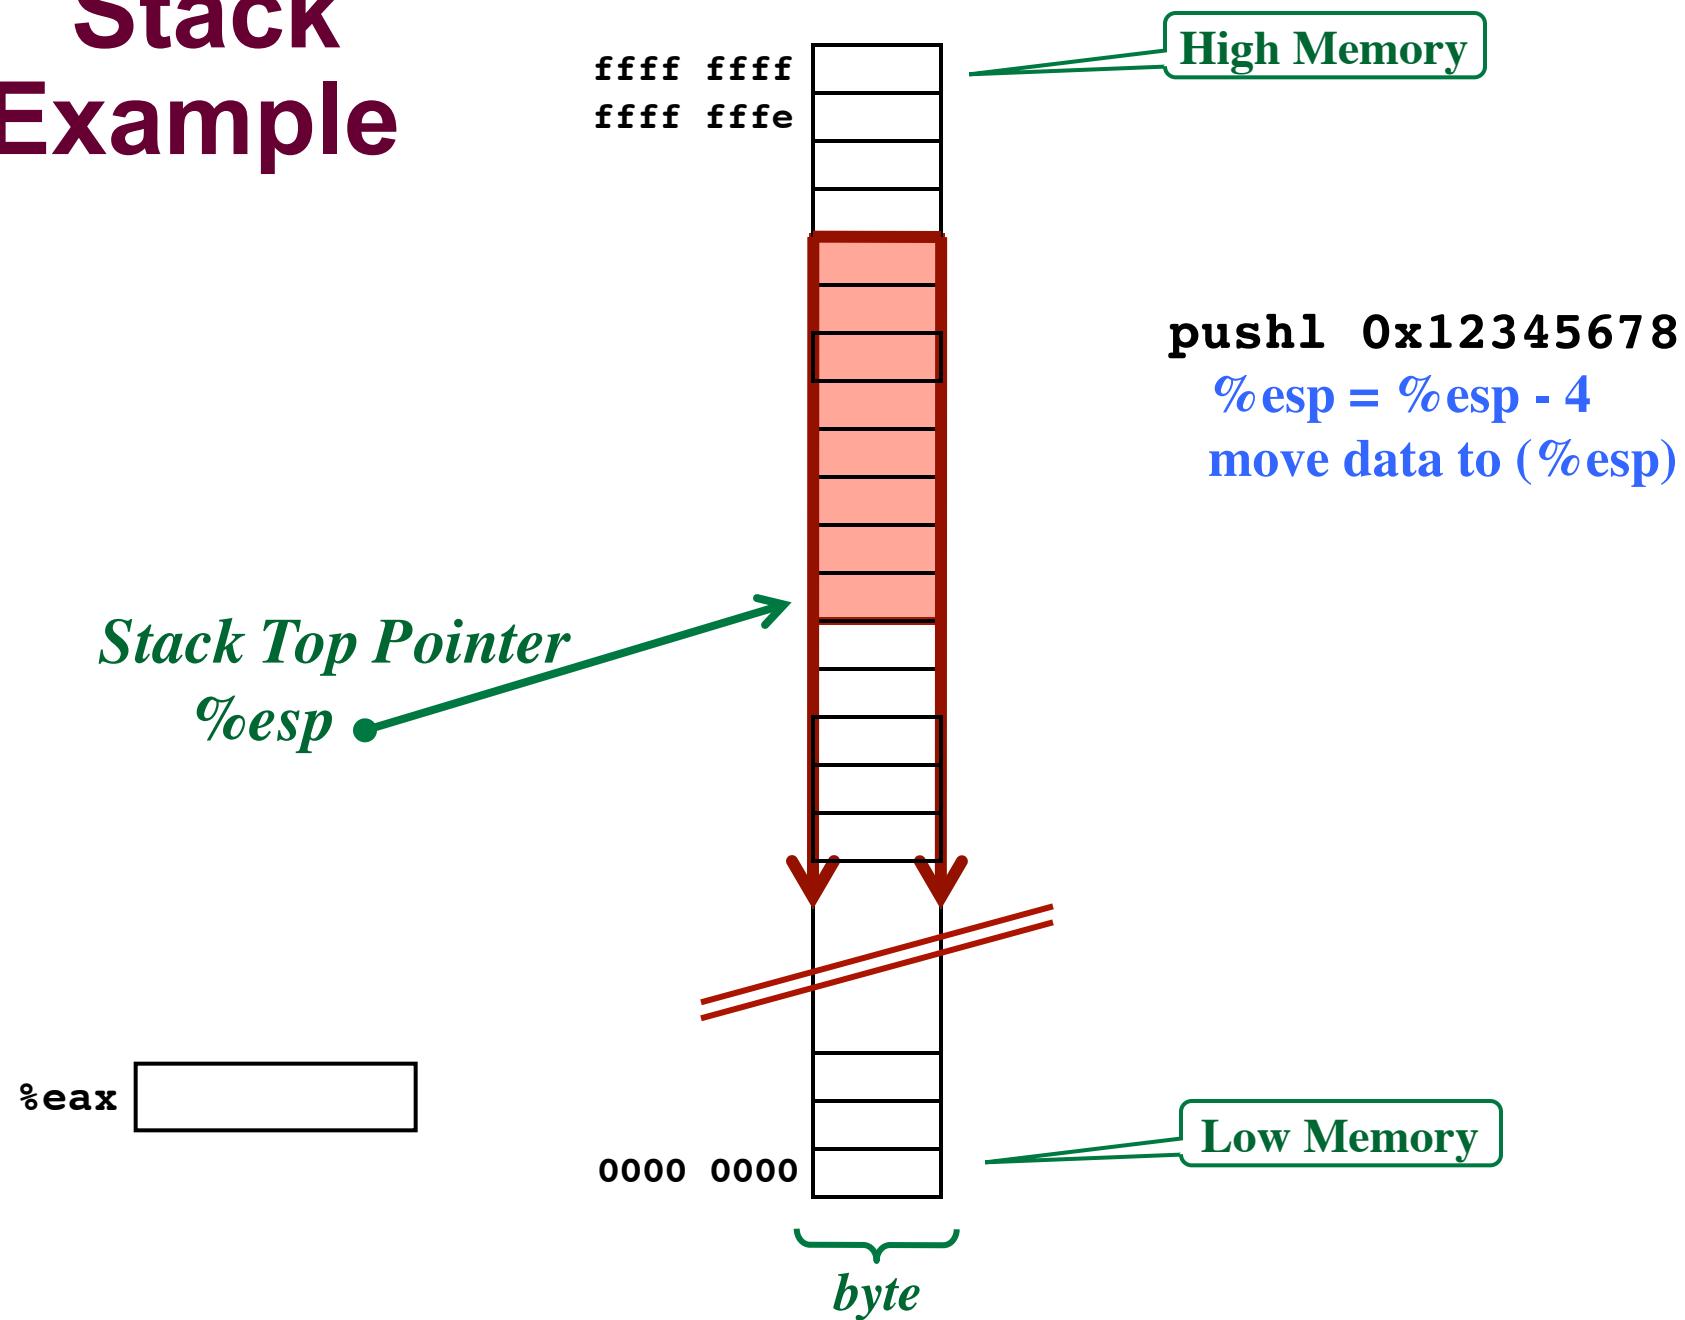
The Stack Grows Downward From High Memory



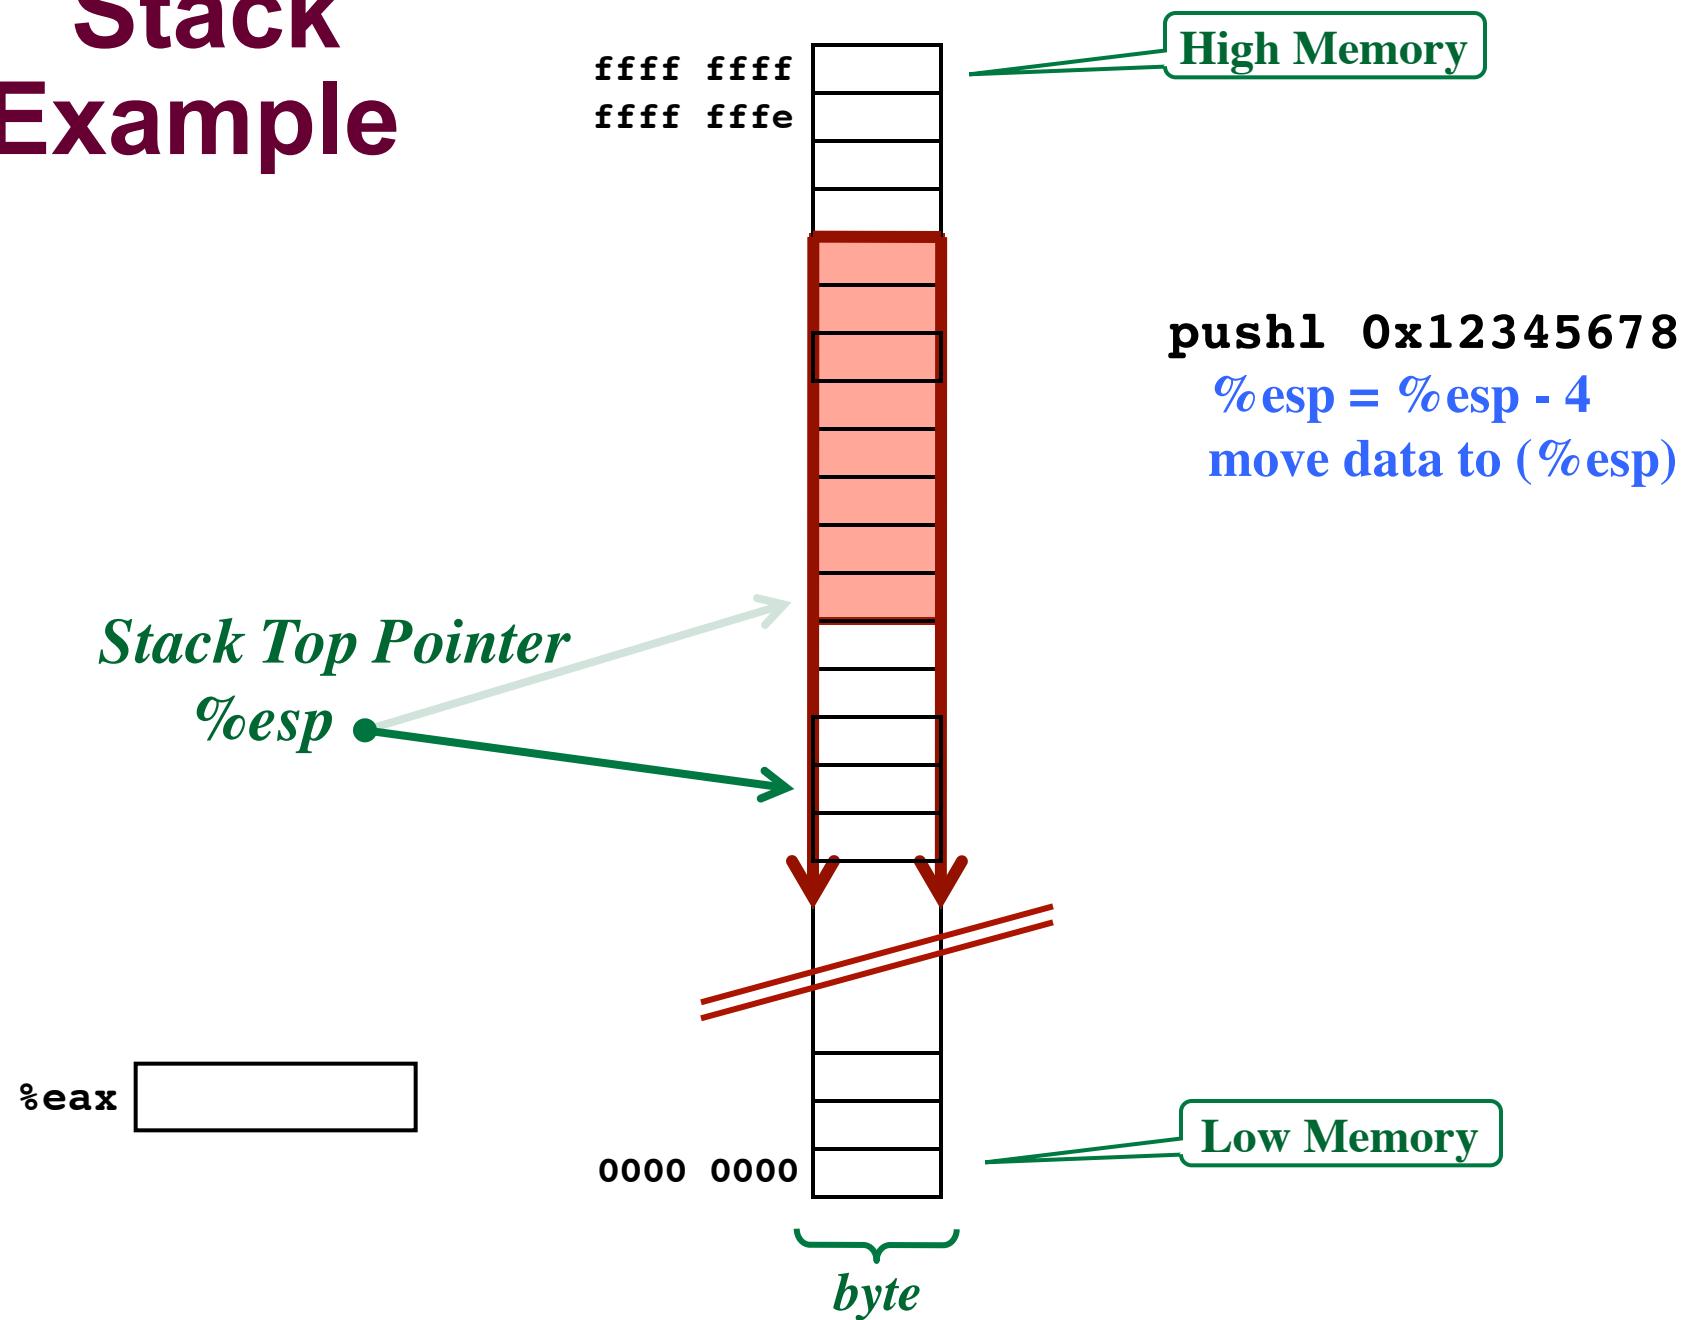
Stack Example



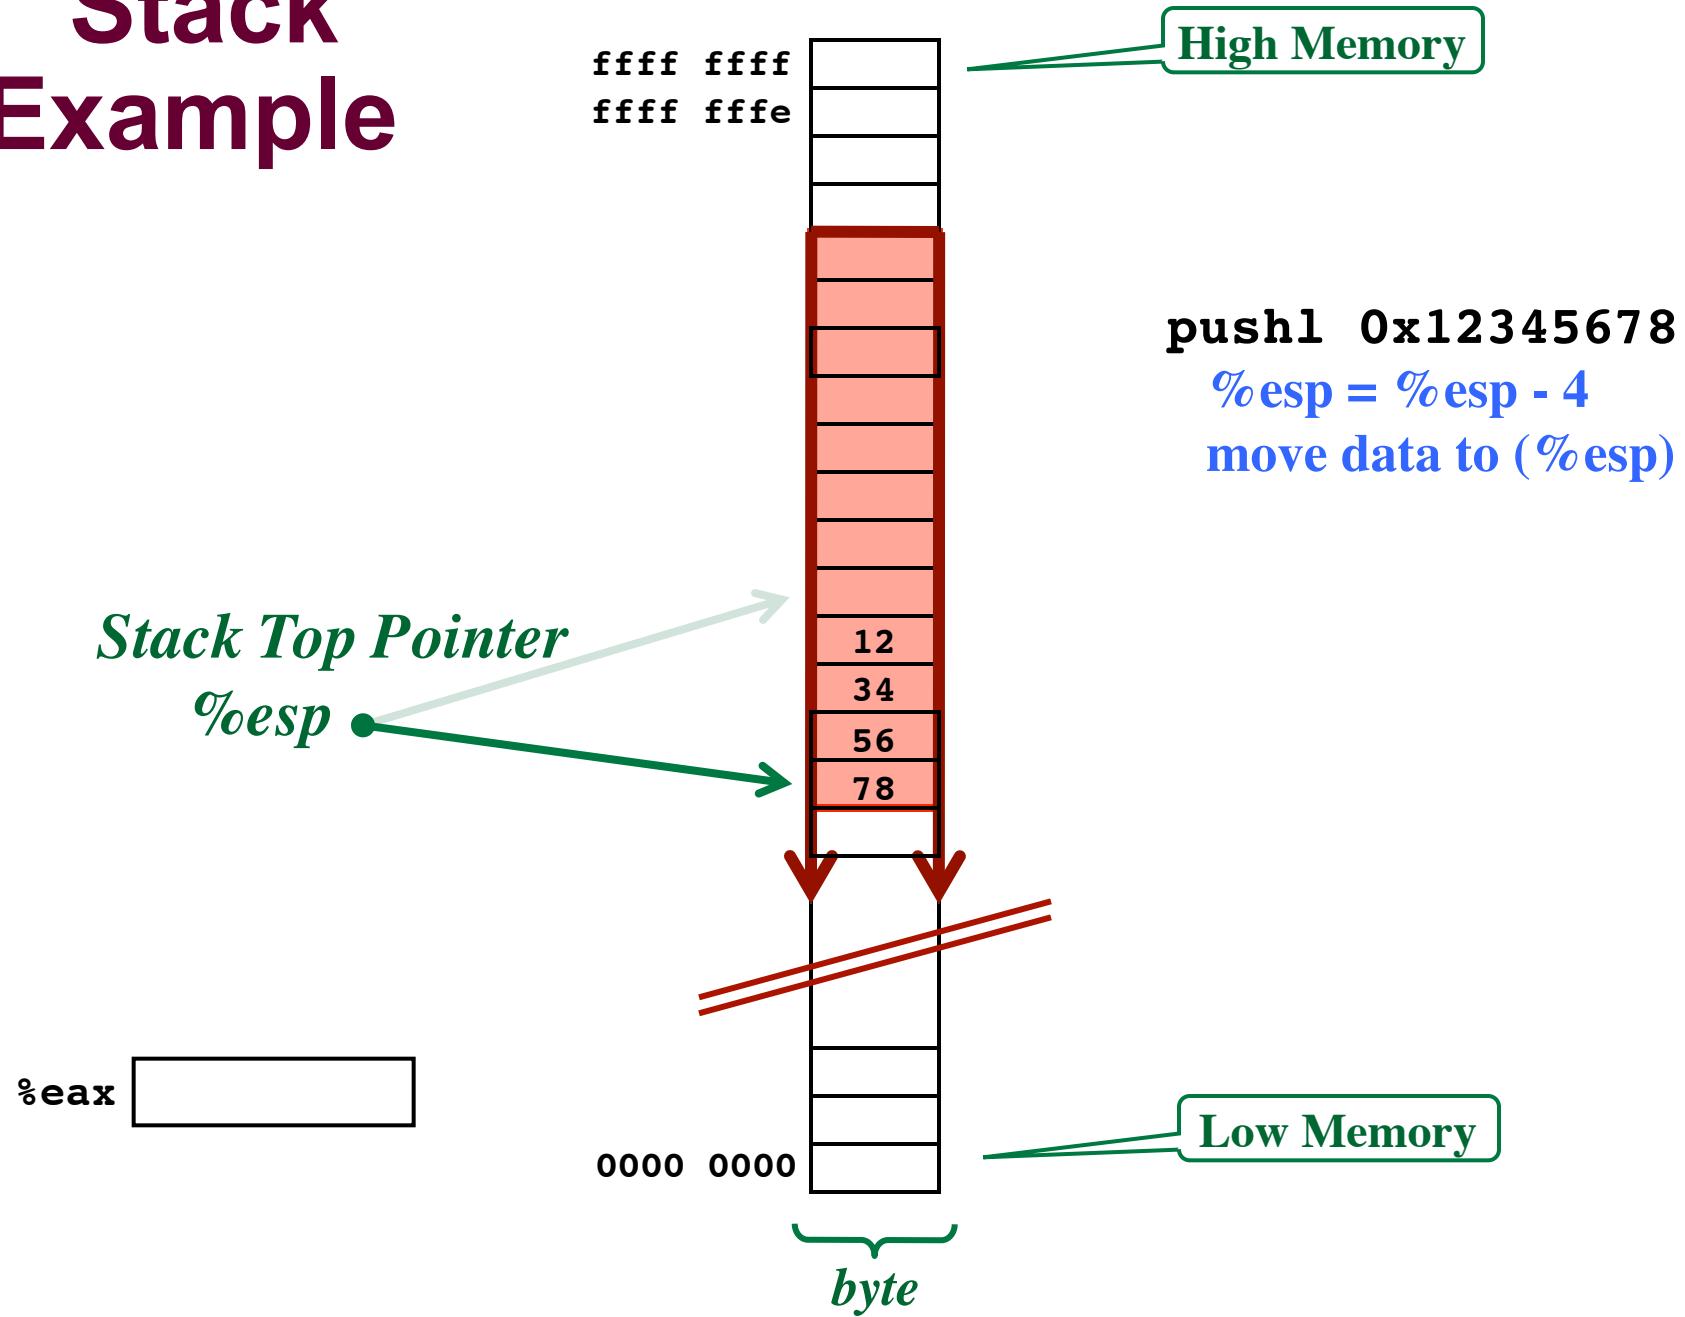
Stack Example



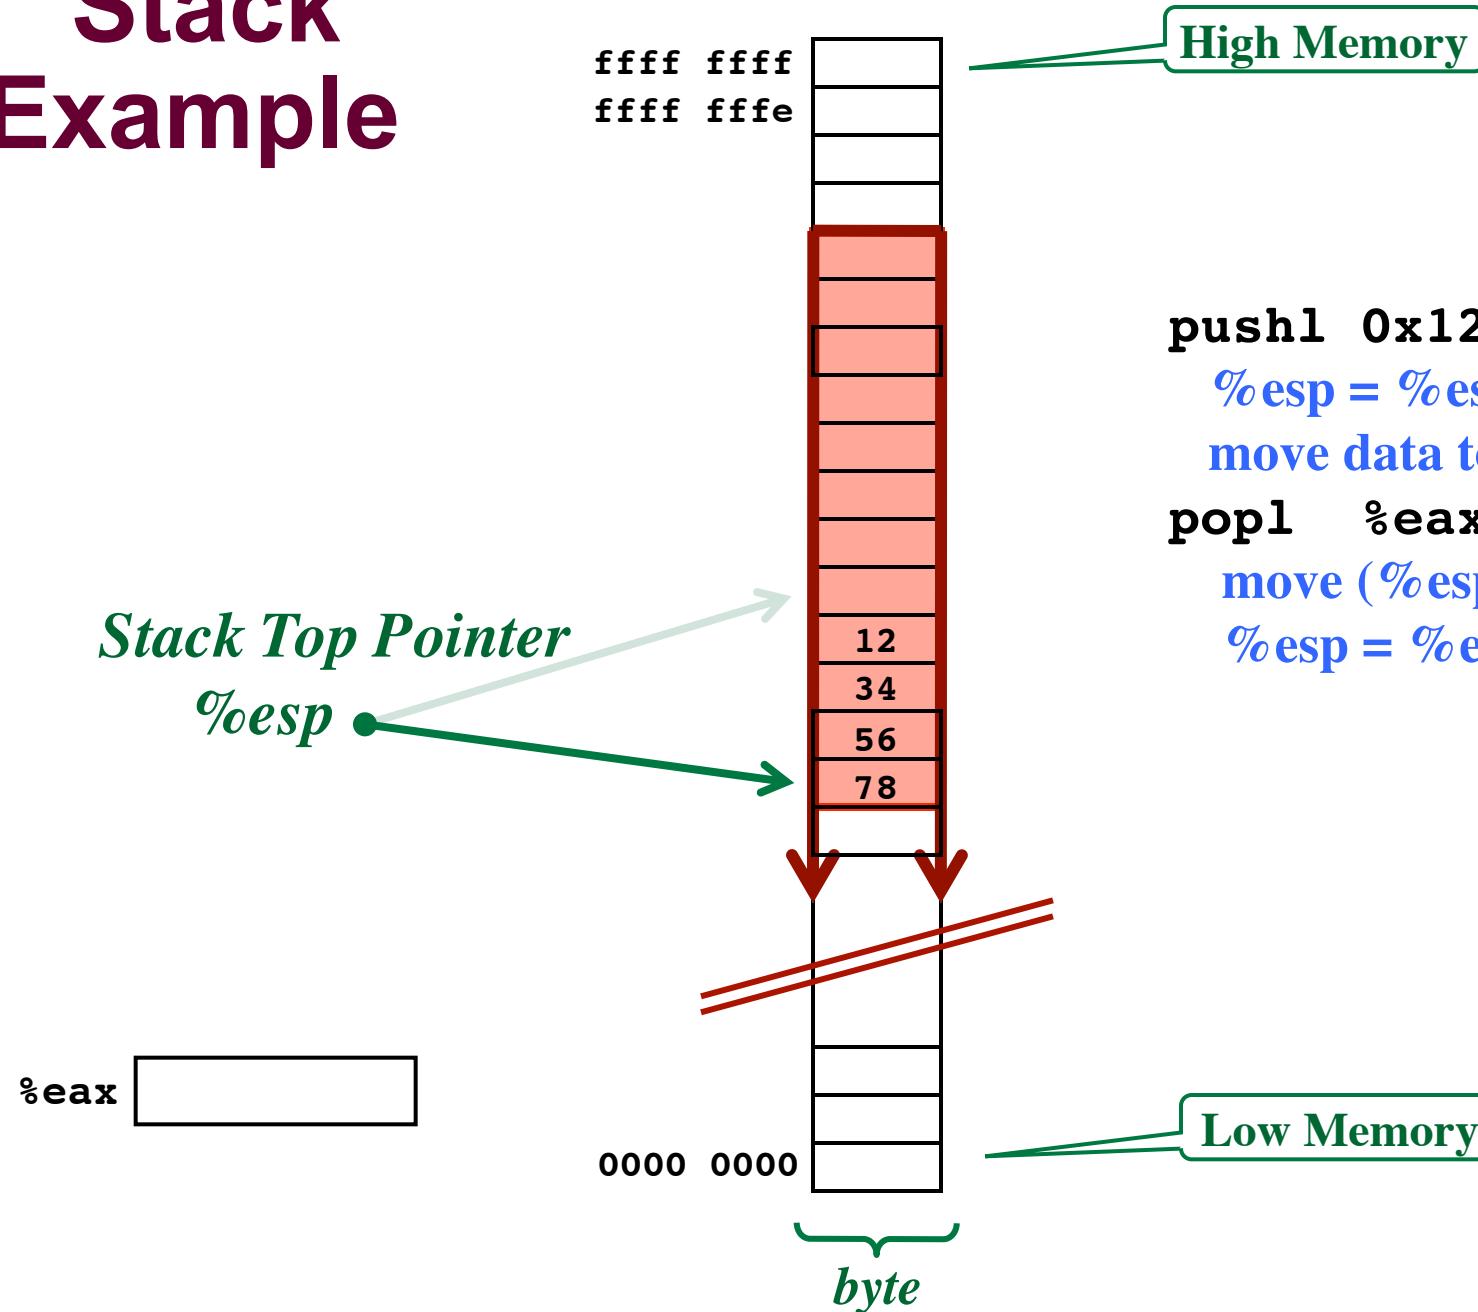
Stack Example



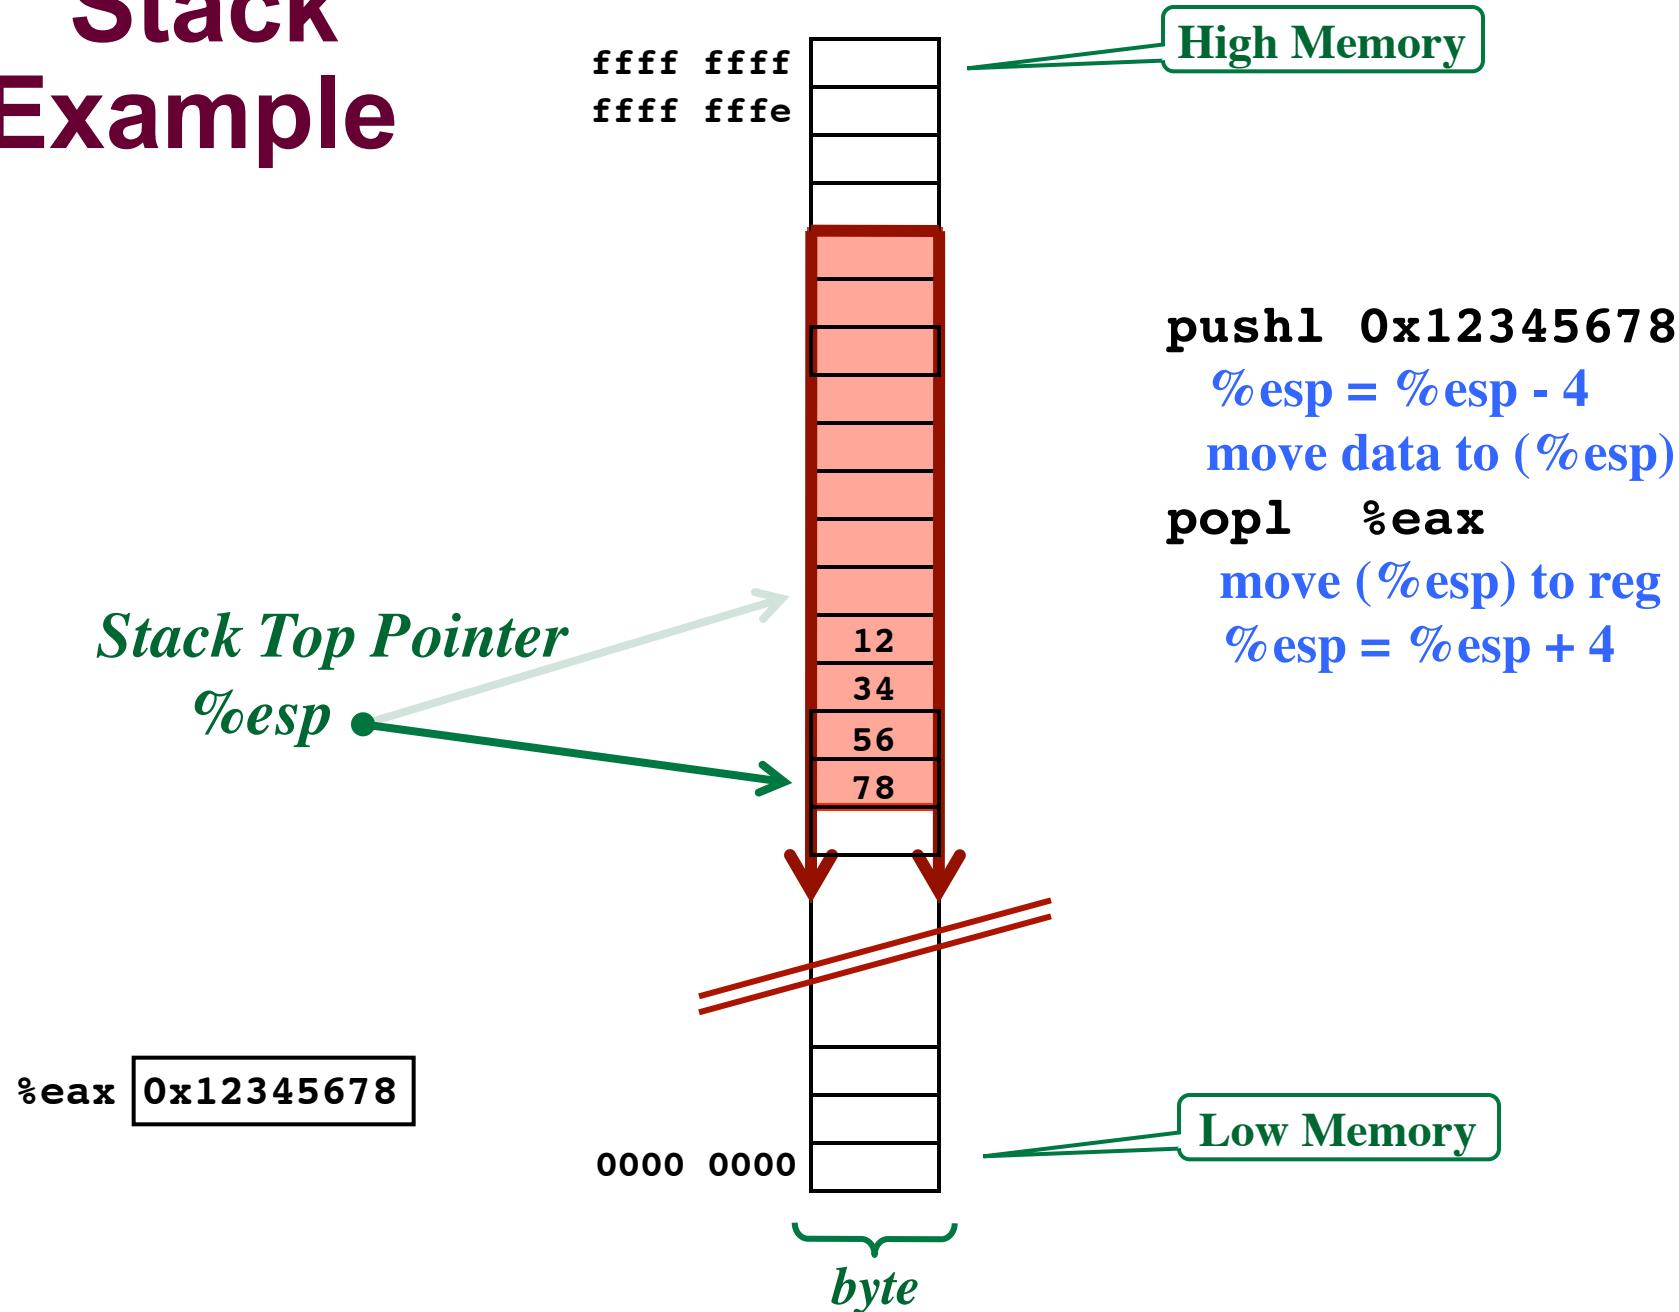
Stack Example



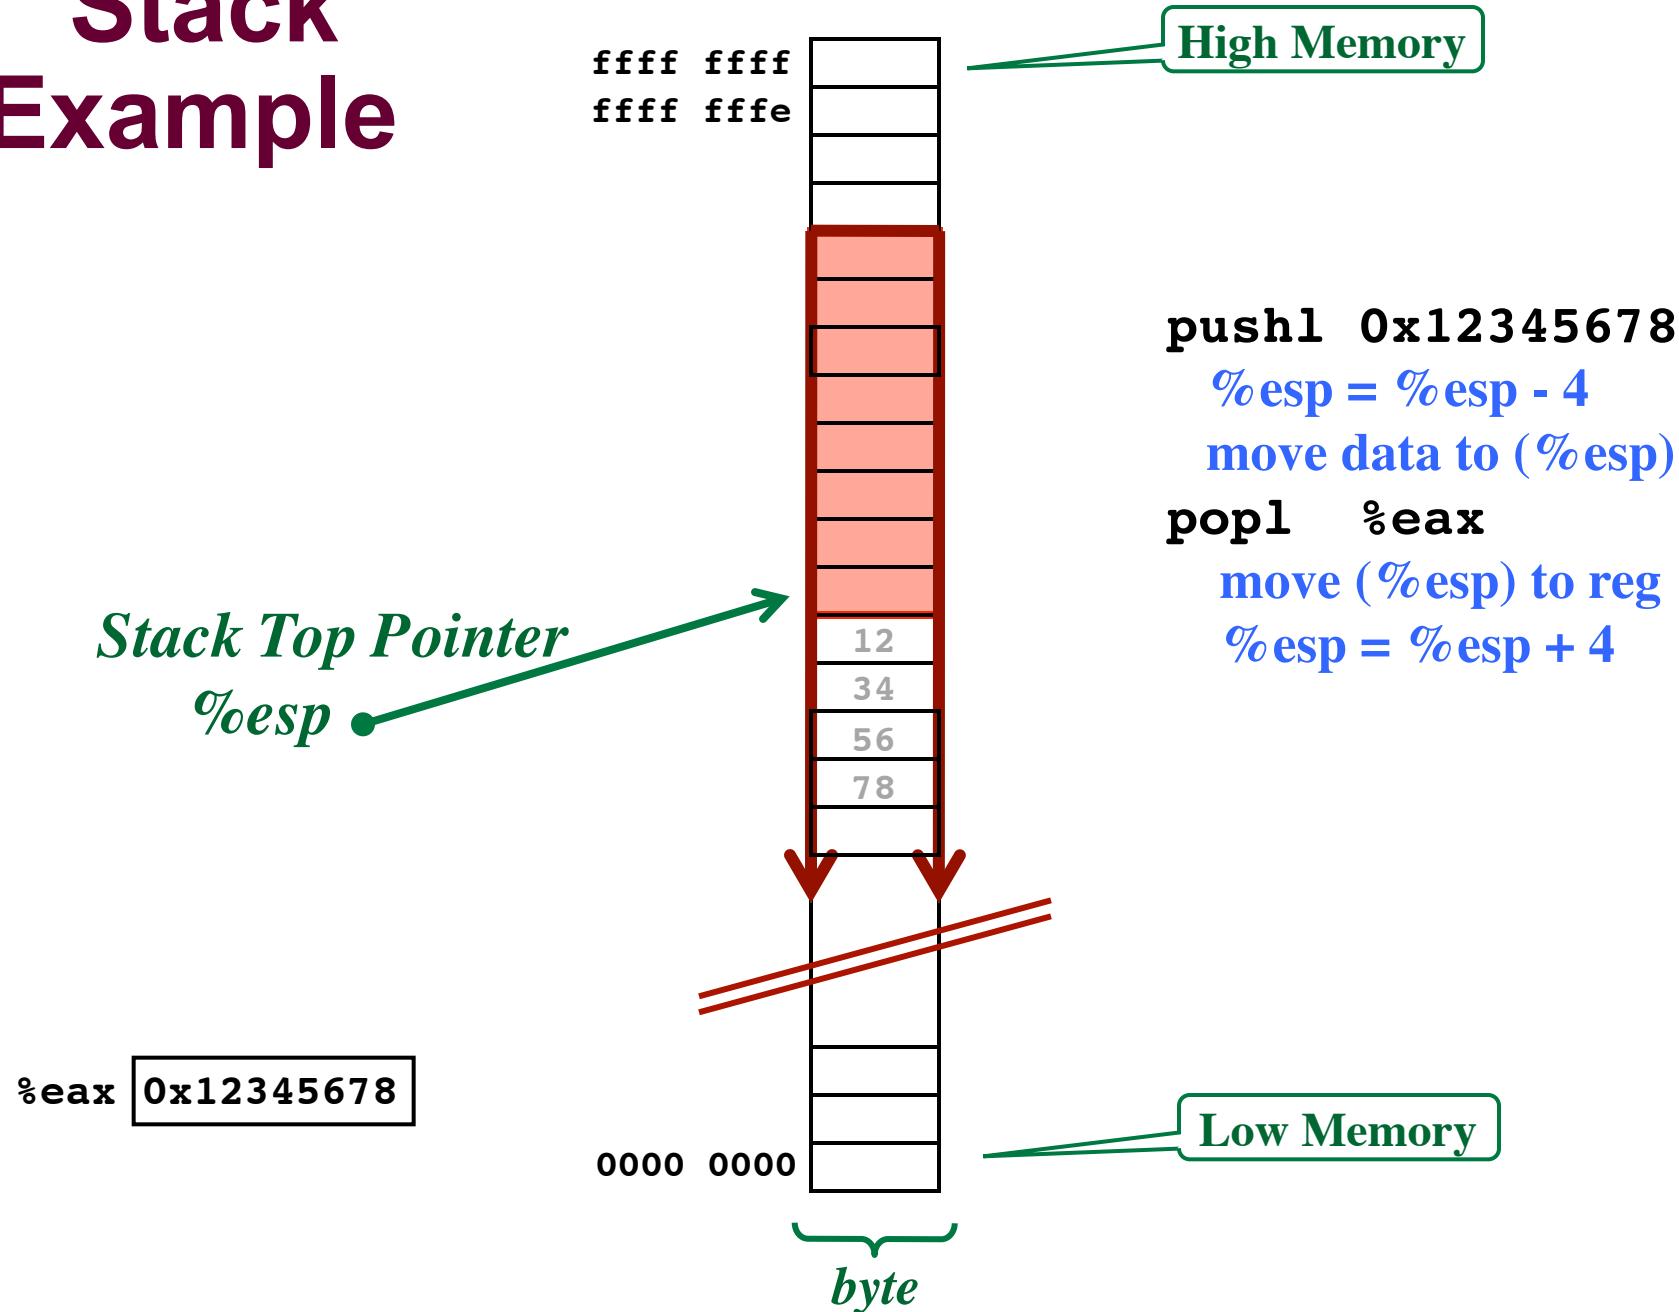
Stack Example



Stack Example



Stack Example



Stack instructions (push and pop)

Stack manipulation is just data movement that updates the stack pointer register

- Move data onto the stack (pushl)
- Move data off of the stack (popl)

The stack is essential for implementing function calls

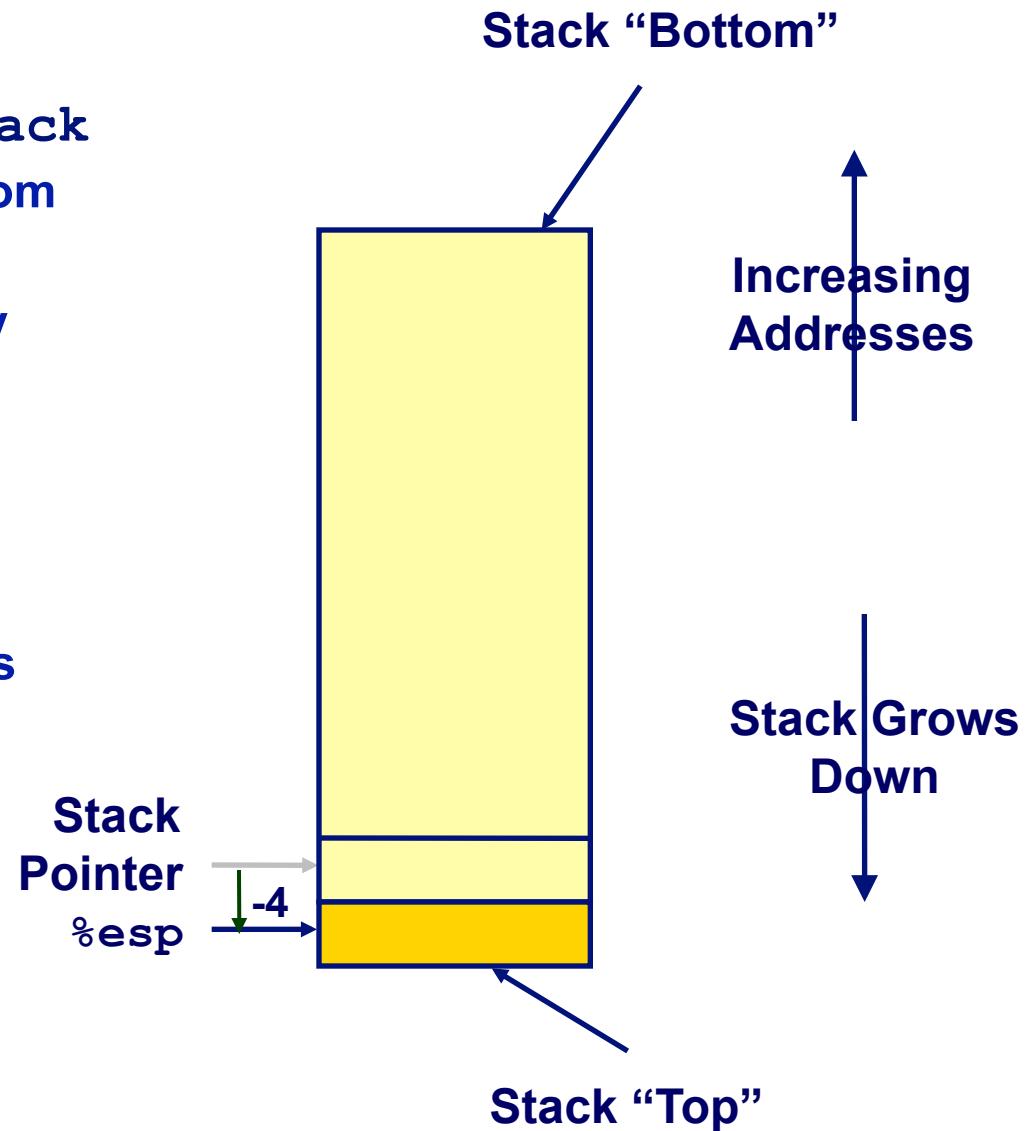
- Function parameters
- Return address
- Prior stack frame information
- Local function variables

Instruction	Effect	Description
<code>pushl</code> S	$R[\%esp] \leftarrow R[\%esp] - 4;$ $M[R[\%esp]] \leftarrow S$	Push
<code>popl</code> D	$D \leftarrow M[R[\%esp]];$ $R[\%esp] \leftarrow R[\%esp] + 4$	Pop

IA32 Stack Pushing

Pushing

- Pushing data onto stack
 - Subtract the data size from the stack pointer (%esp)
 - Move data into the newly created space
- **pushl Src**
 - Fetch operand at *Src*
 - Decrement %esp by 4
 - Write operand at address given by %esp
- **e.g. pushl %eax**
`subl $4, %esp`
`movl %eax, (%esp)`

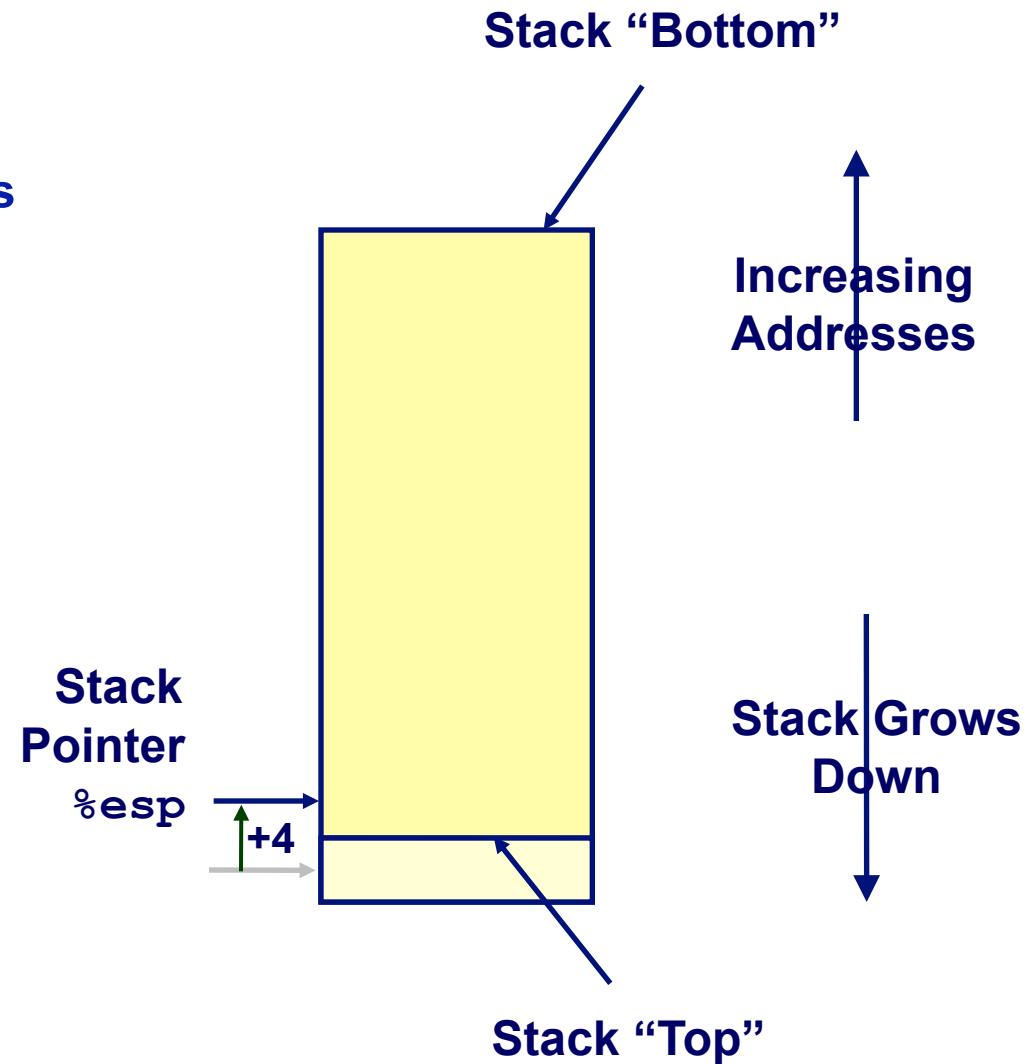


IA32 Stack Popping

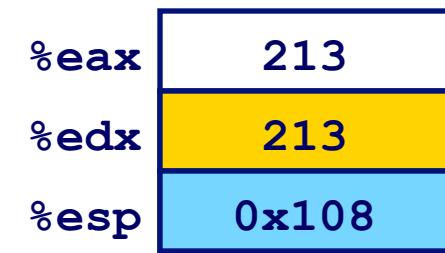
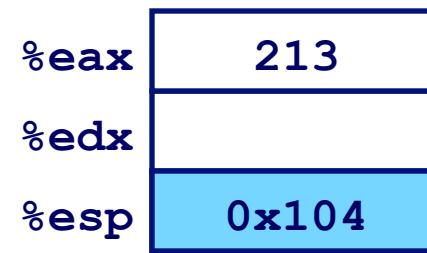
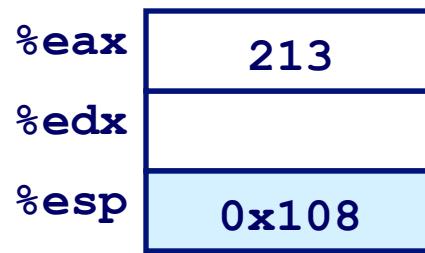
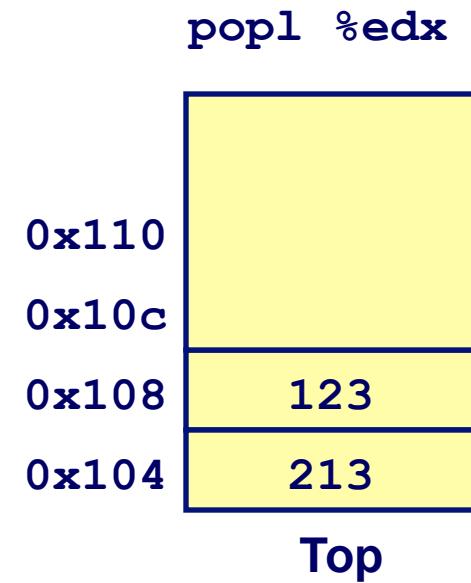
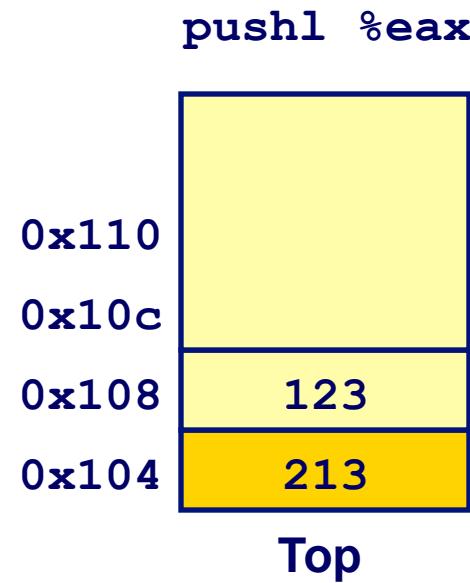
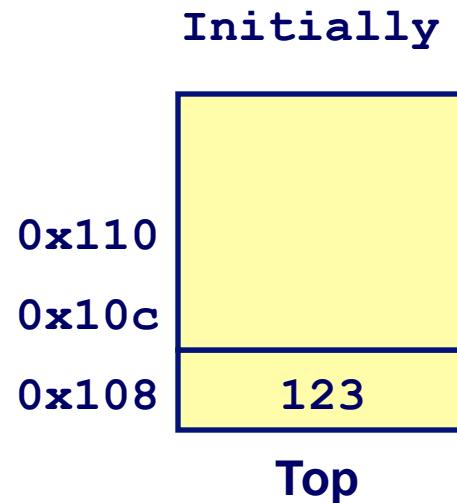
Popping

- **popl Dest**
 - Read operand at address given by `%esp`
 - Write to *Dest*
 - Increment `%esp` by 4
- **e.g. popl %eax**

```
movl (%esp), %eax
addl $4, %esp
```



Stack Operation Examples



Procedure Control Flow

Procedure call:

`call label`

- Push address of next instruction (after the call) on stack
- This is your return address
- Jump to `label`

Procedure return:

- `ret` Pop return address from stack into eip register

Procedure Control Flow

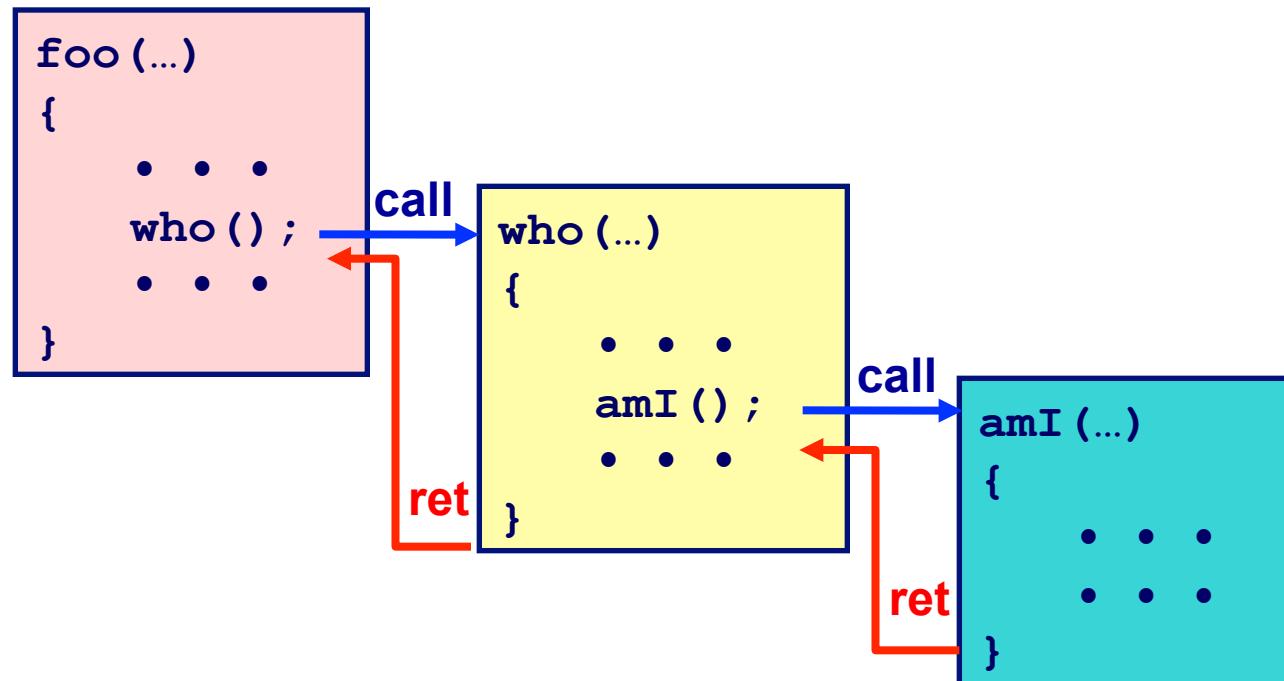
When procedure `foo` calls `who`:

- `foo` is the *caller*, `who` is the *callee*
- Control is transferred to the ‘callee’

When procedure returns

- Control is transferred back to the ‘caller’

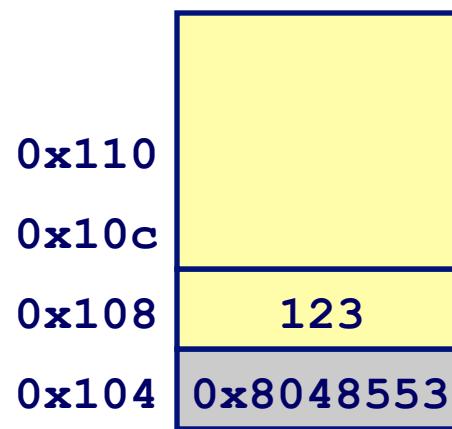
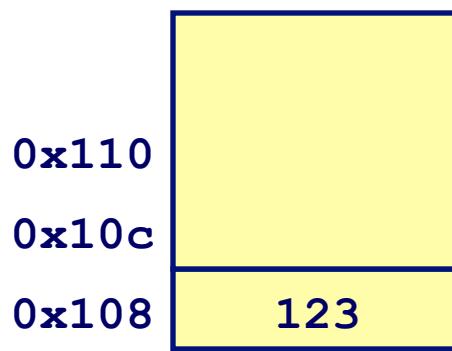
Last-called, first-return (LIFO) order naturally implemented via stack



Procedure Call Example

```
804854e: e8 3d 06 00 00      call    8048b90 <main>
8048553: 50                  next instruction
```

call 8048b90



%esp 0x108
%eip 0x804854e

%esp 0x104
%eip 0x8048b90

%eip *is program counter*

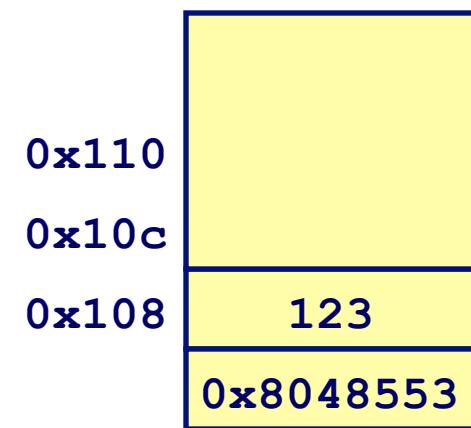
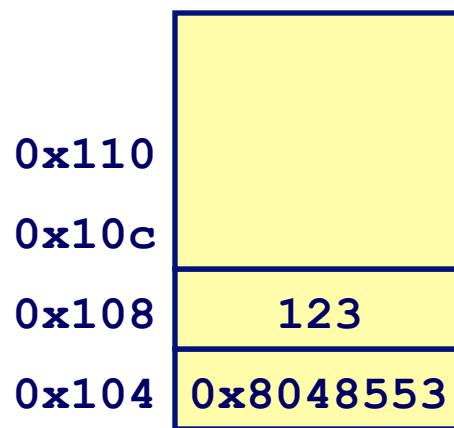
Procedure Return Example

animation

8048e90: c3

ret

ret



%esp | 0x104

%esp 0x108

%eip 0x8048e90

%eip 0x8048553

%eip *is program counter*

Keeping track of stack frames

The stack pointer (%esp) moves around

Can be changed within a procedure

Problem

How can we consistently find our local variables, return address and parameters?

The base pointer (%ebp)

Points to the base of our current stack frame

Also called the frame pointer

Within each function, %ebp stays constant

Most information on the stack is referenced relative to the base pointer

- Base pointer setup is the programmer's job
Actually usually the compiler's job

Procedure calls and stack frames

How does the ‘callee’ know where to return later?

- Return address placed in a well-known location on stack within a “stack frame”

How are arguments passed to the ‘callee’?

- Arguments placed in a well-known location on stack within a “stack frame”

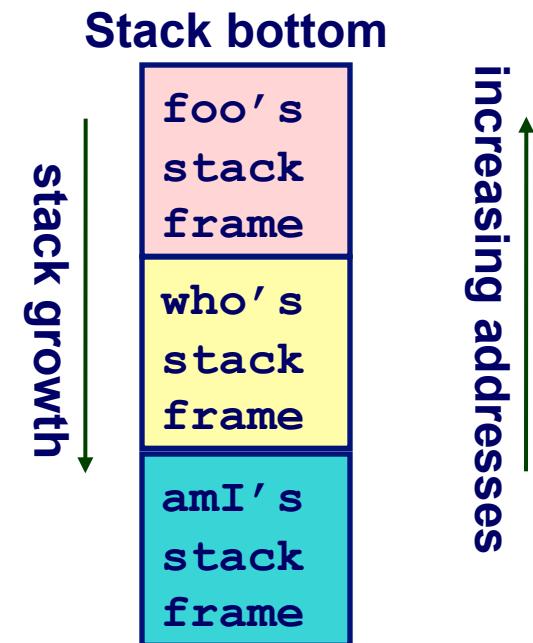
Upon procedure invocation

- Stack frame created for the procedure
- Stack frame is pushed onto program stack

Upon procedure return

- Its frame is popped off of stack
- Caller’s stack frame is recovered

Call chain: foo => who => amI



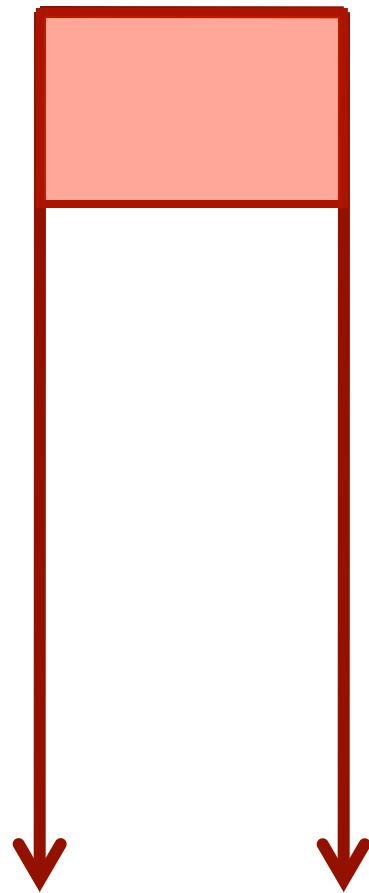
Functions and Stack Frames

```
main() {  
    ... call f1()  
}
```

```
f1() {  
    ... call f2()  
}
```

```
f2() {  
    ... call f3()  
}
```

```
f3 () {  
    ...  
}
```



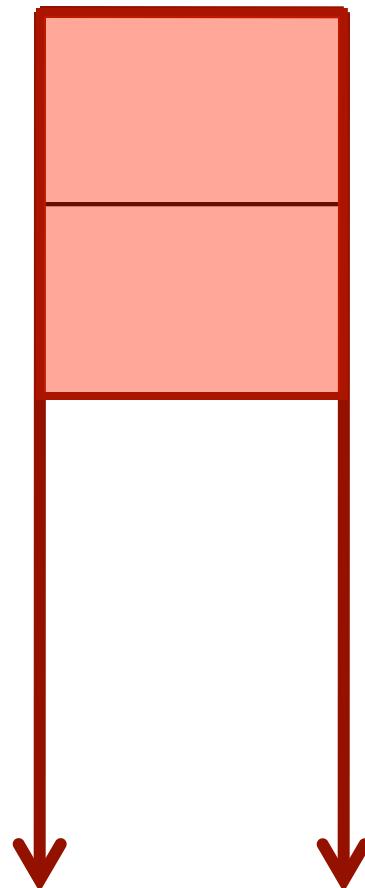
Functions and Stack Frames

```
main() {  
    ... call f1()  
}
```

```
f1() {  
    ... call f2()  
}
```

```
f2() {  
    ... call f3()  
}
```

```
f3 () {  
    ...  
}
```



Frame of main

Frame of f1

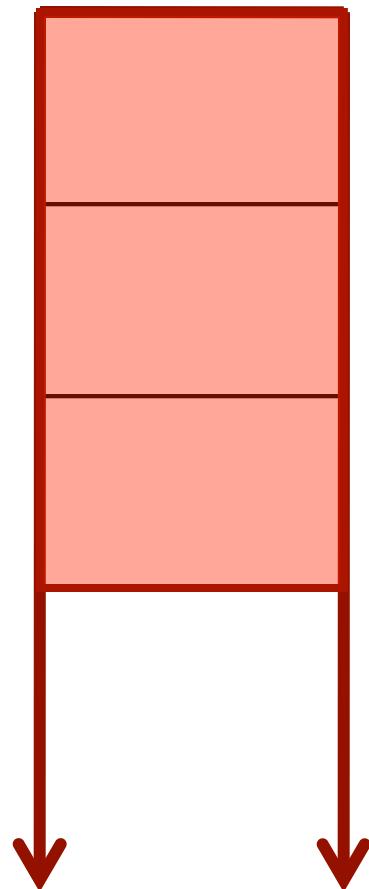
Functions and Stack Frames

```
main() {  
    ... call f1()  
}
```

```
f1() {  
    ... call f2()  
}
```

```
f2() {  
    ... call f3()  
}
```

```
f3 () {  
    ...  
}
```



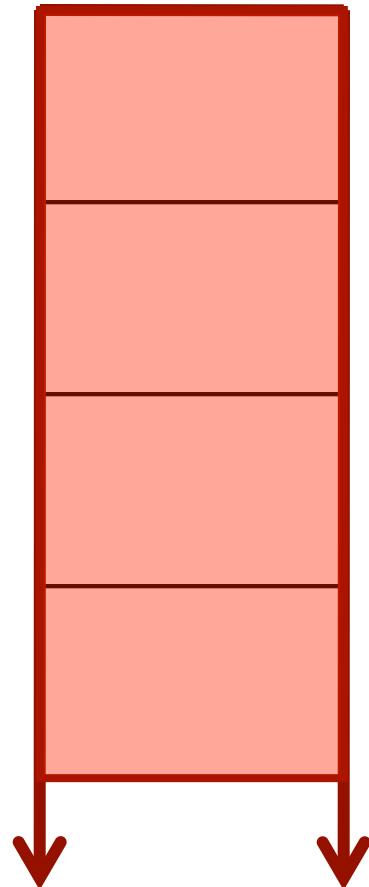
Frame of main

Frame of f1

Frame of f2

Functions and Stack Frames

```
main() {  
    ... call f1()  
}  
  
f1() {  
    ... call f2()  
}  
  
f2() {  
    ... call f3()  
}  
  
f3 () {  
    ...  
}
```



Frame of main

Frame of f1

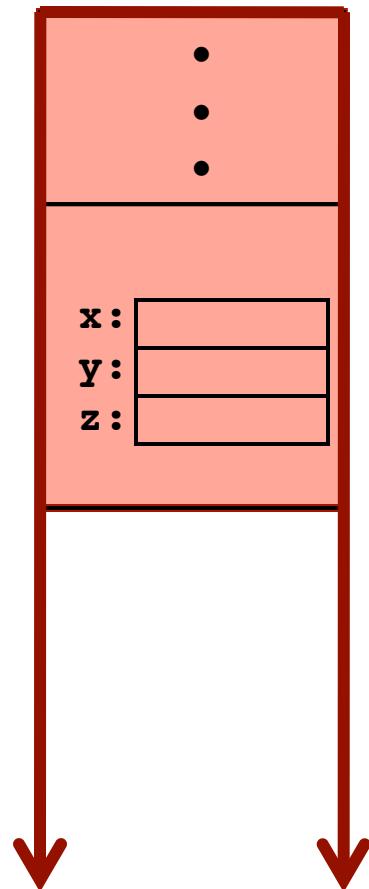
Frame of f2

Frame of f3

Functions and Stack Frames

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

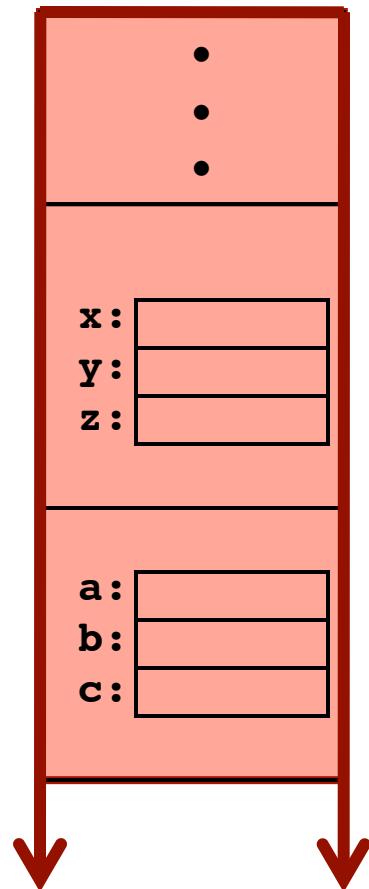


Frame of f1

Functions and Stack Frames

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```



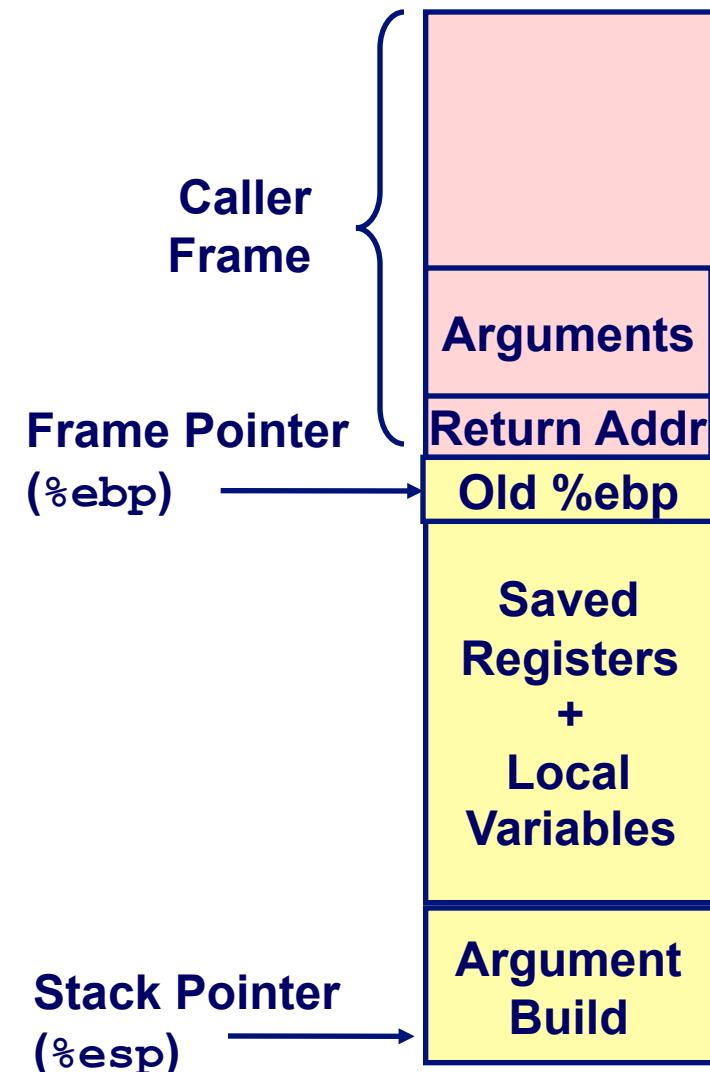
IA32/Linux Stack Frame

Caller Stack Frame (Pink)

- Argument build for callee
 - i.e. function parameters of callee
- Return address
 - Pushed by `call` instruction

Callee Stack Frame (Yellow) (From Top to Bottom)

- Old frame pointer
- Local variables
 - If can't keep in registers
- Saved register context
- Argument build for next call
 - i.e. function parameters for next call



Stack Frame Creation and Destruction

```
f2 (...) {  
    ...  
}
```

f2:

pushl **%ebp**

movl **%esp, %ebp**

...

movl **%ebp, %esp**

popl **%ebp**

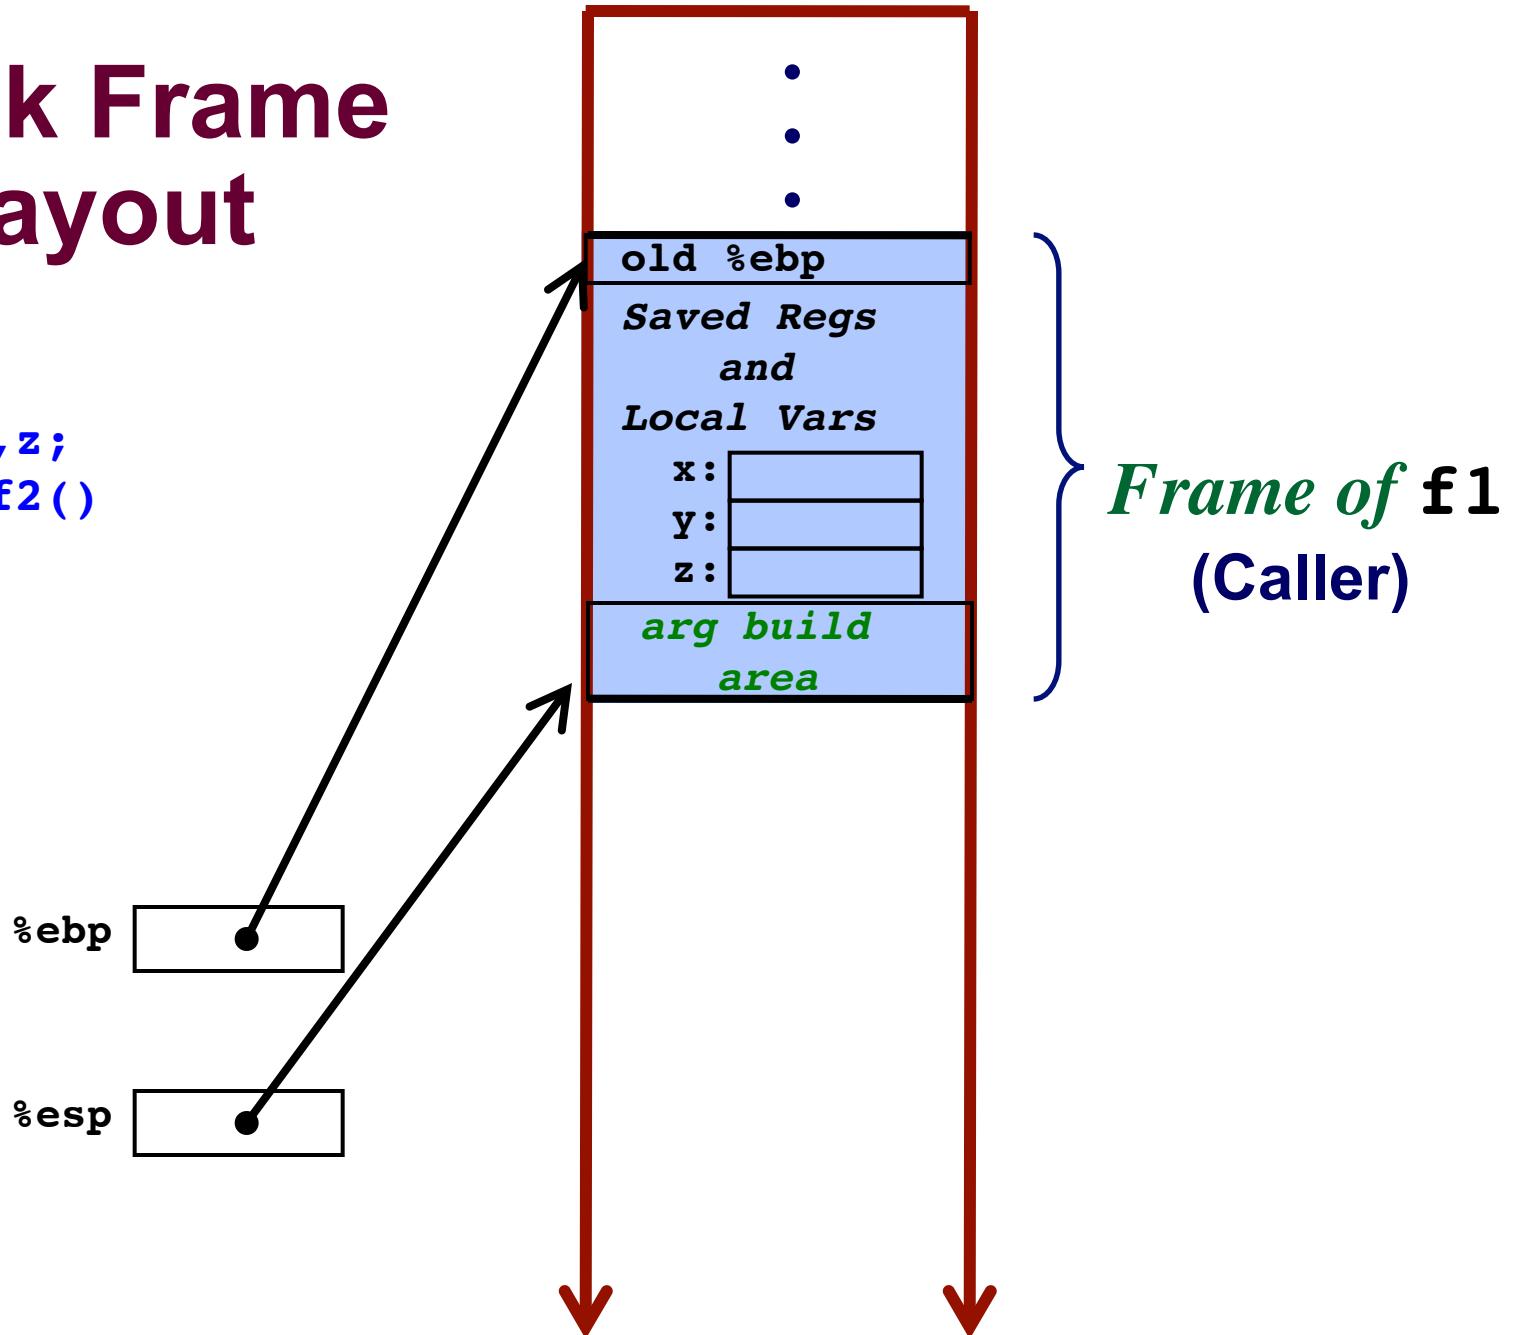
ret

} Prologue

} Epilogue

Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```



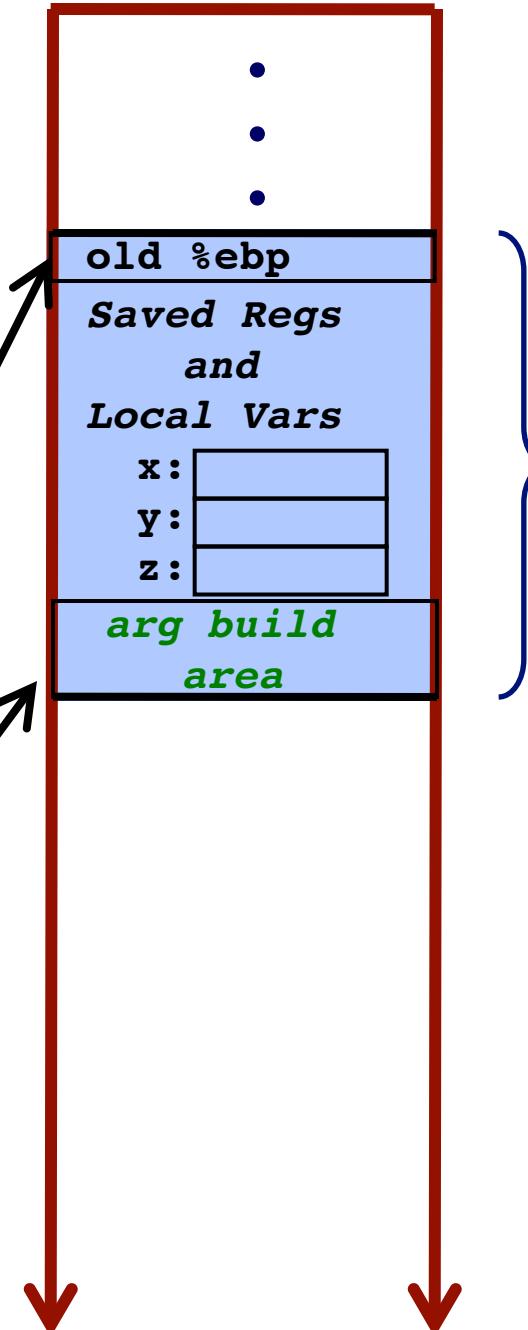
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp



Stack Frame Layout

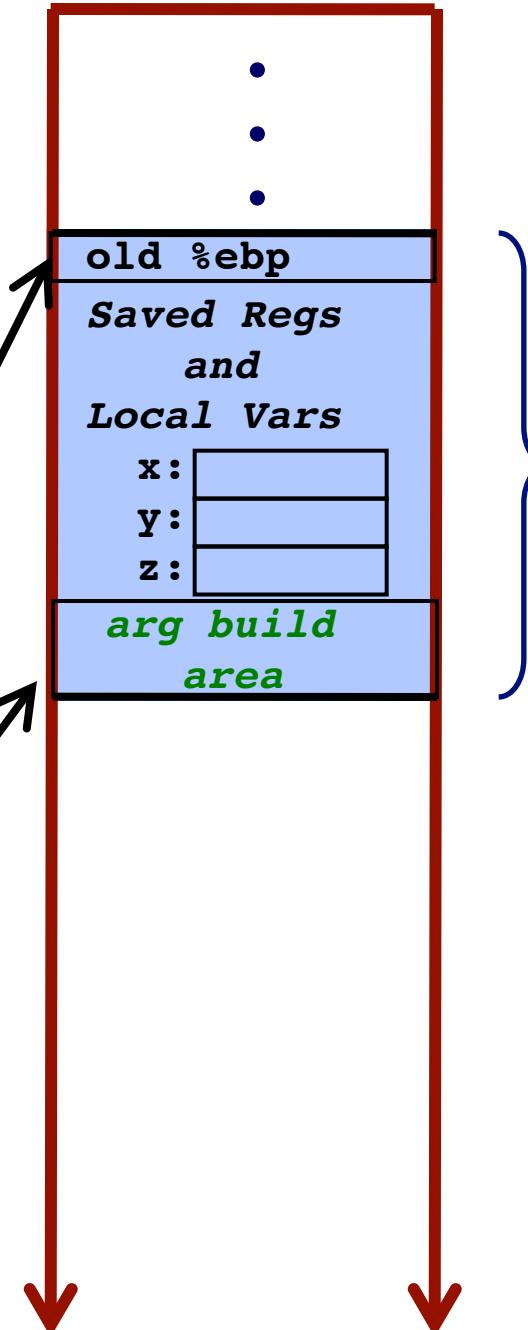
```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp

Compute args to f2 and move into “arg build area”

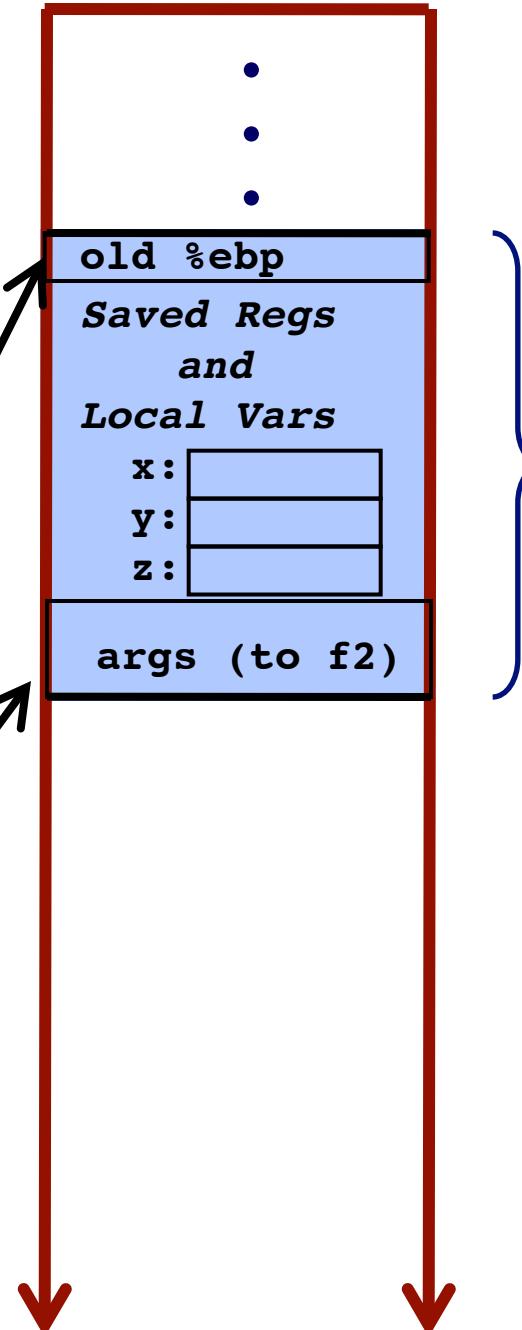
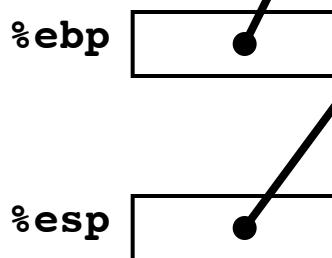


Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

next instruction:
call f2



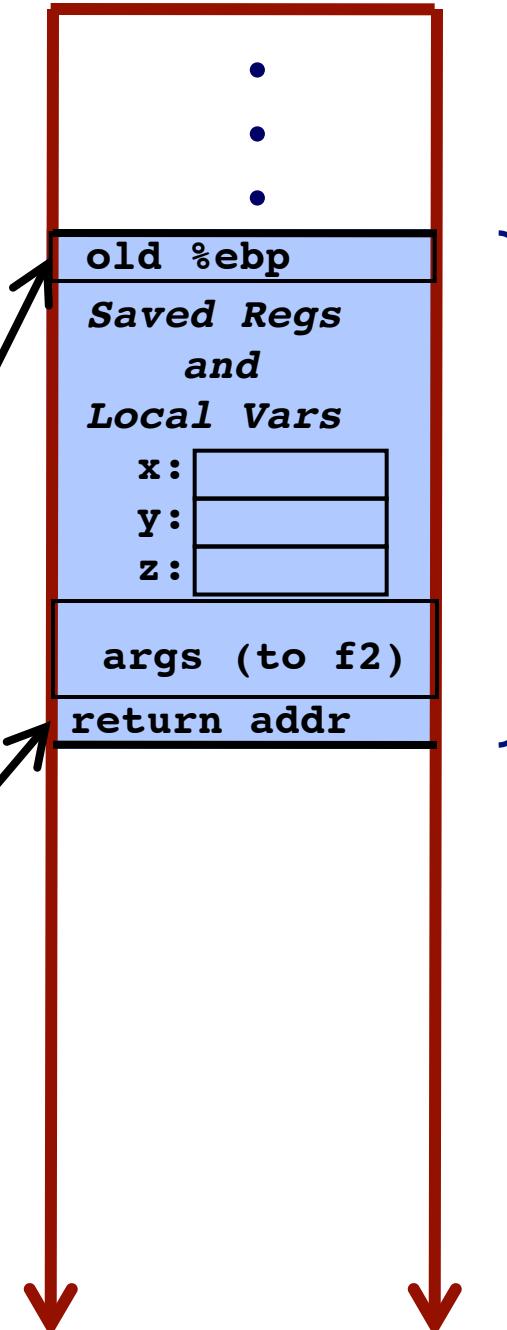
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp



Stack Frame Layout

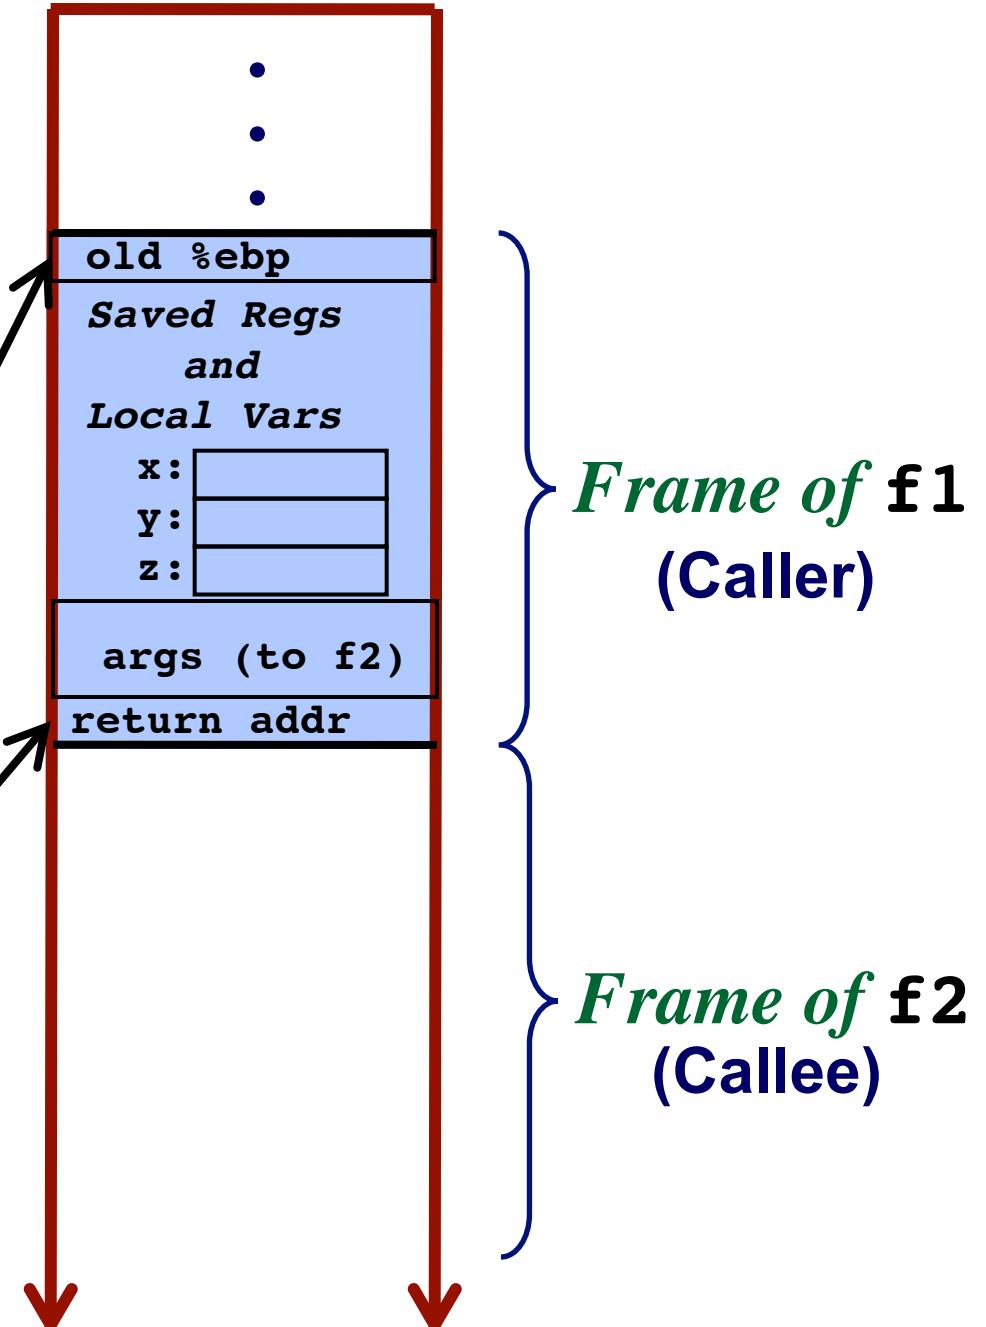
```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp

next instruction:
pushl %ebp



Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

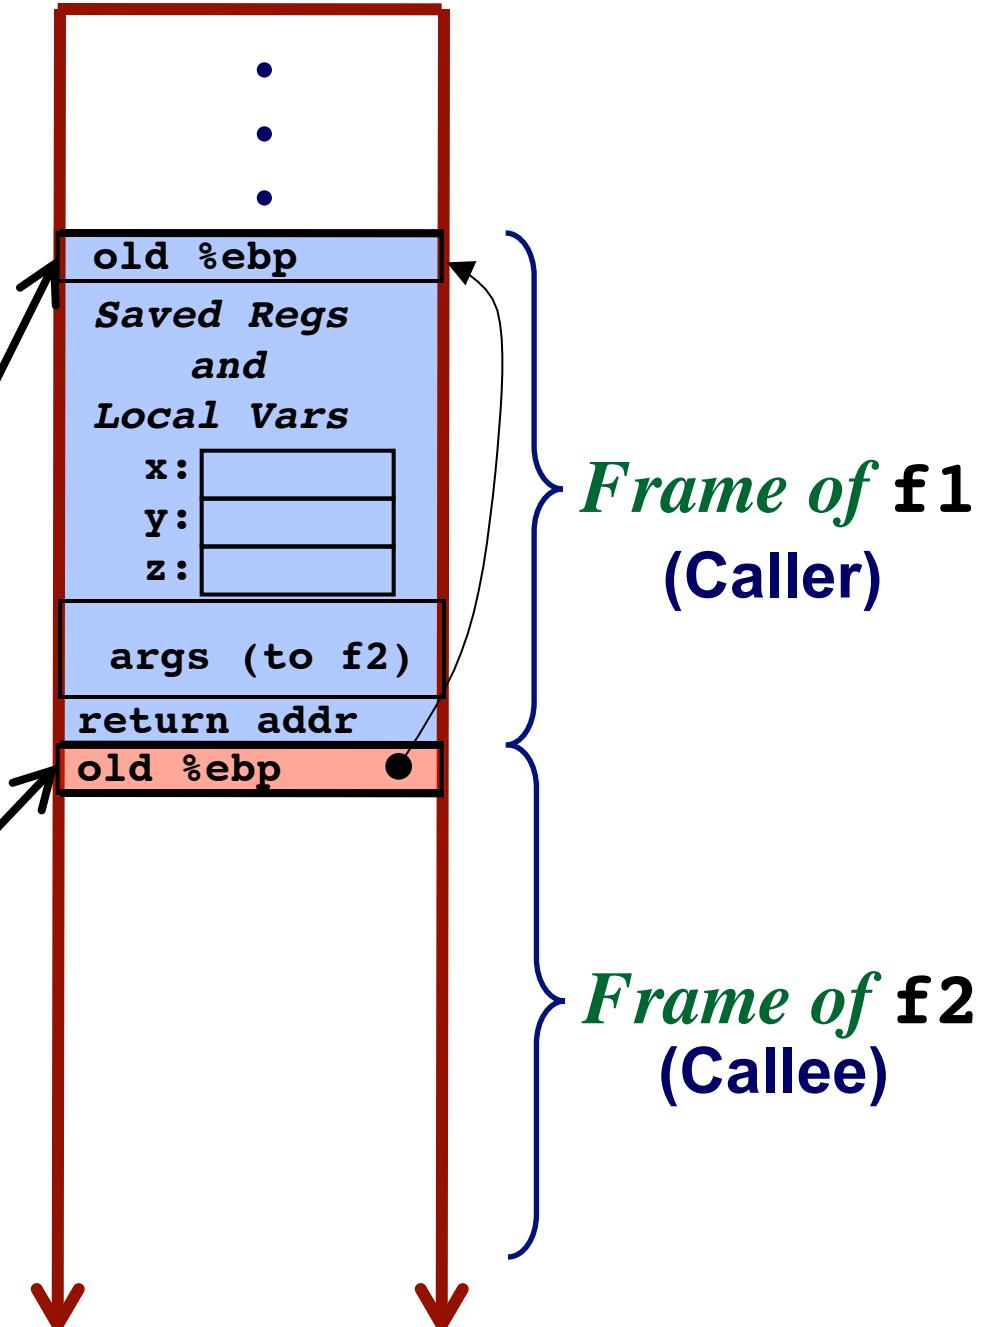
```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp

next instruction:

```
movl %esp,%ebp
```

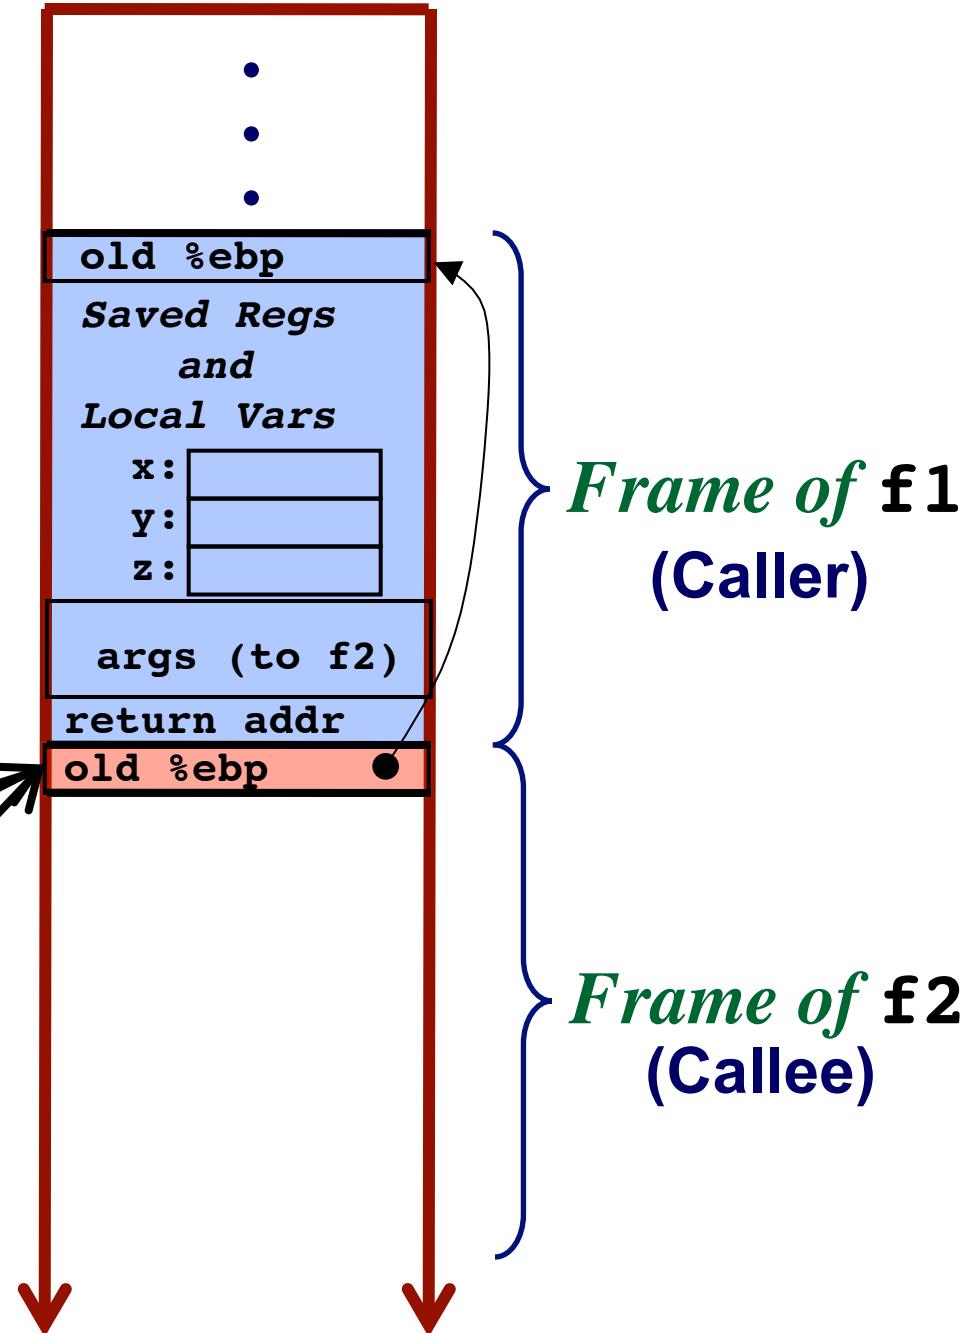


Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp
%esp



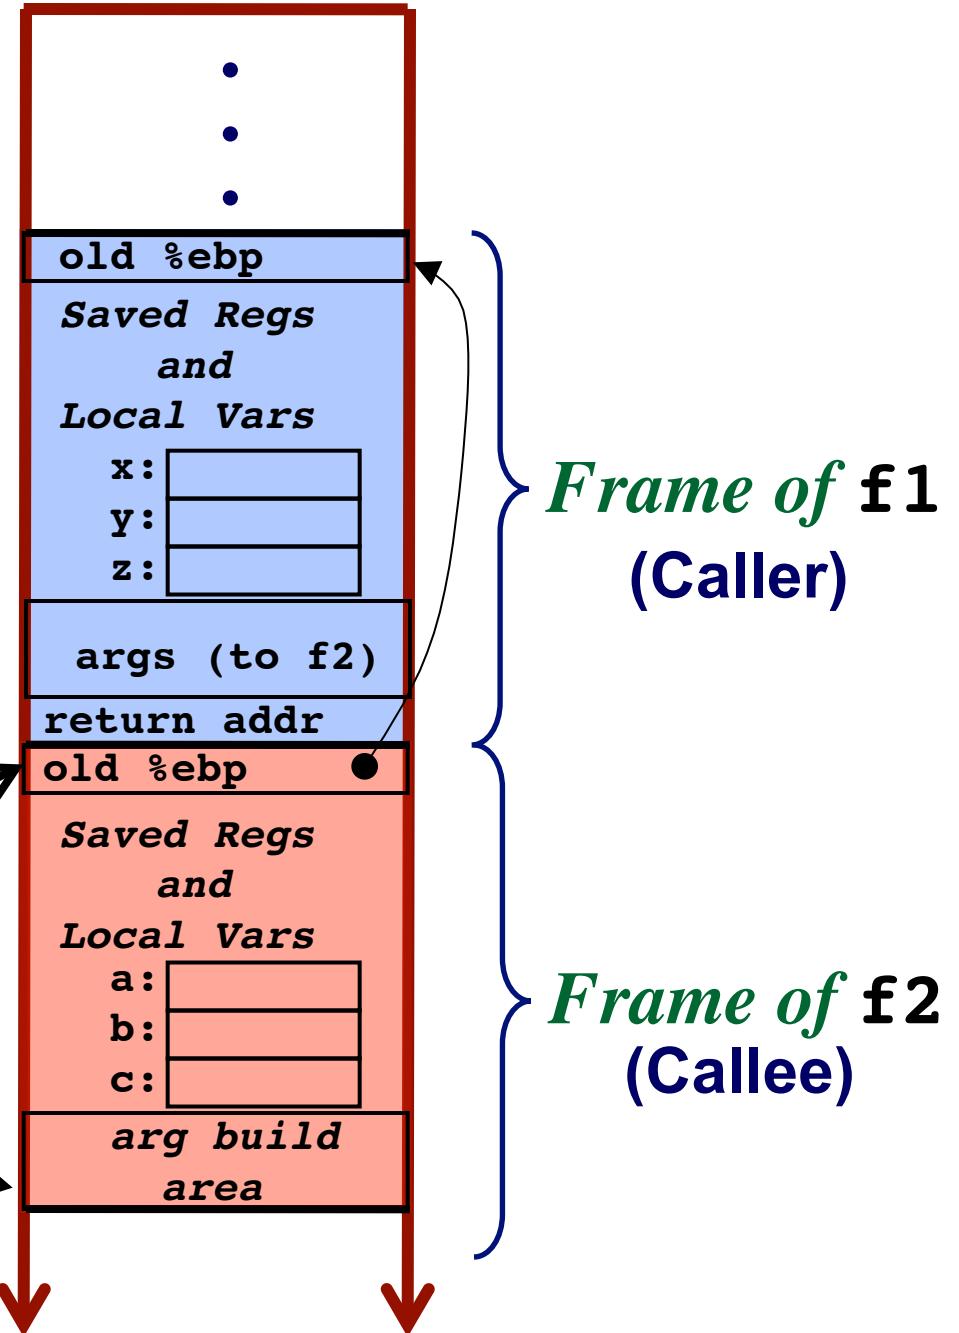
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp



Stack Frame Layout

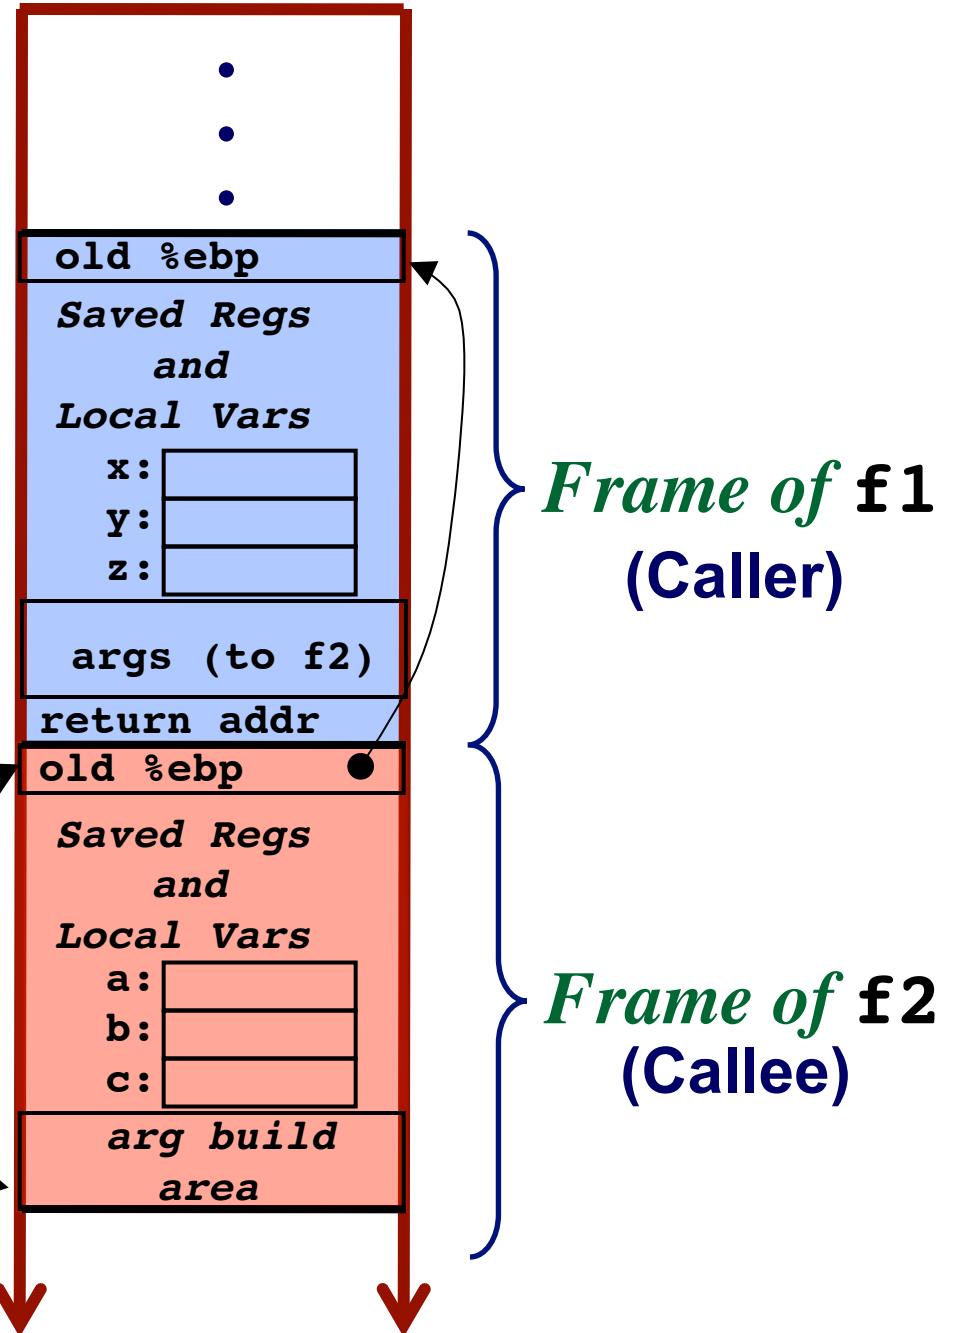
```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp

Want to call f3?



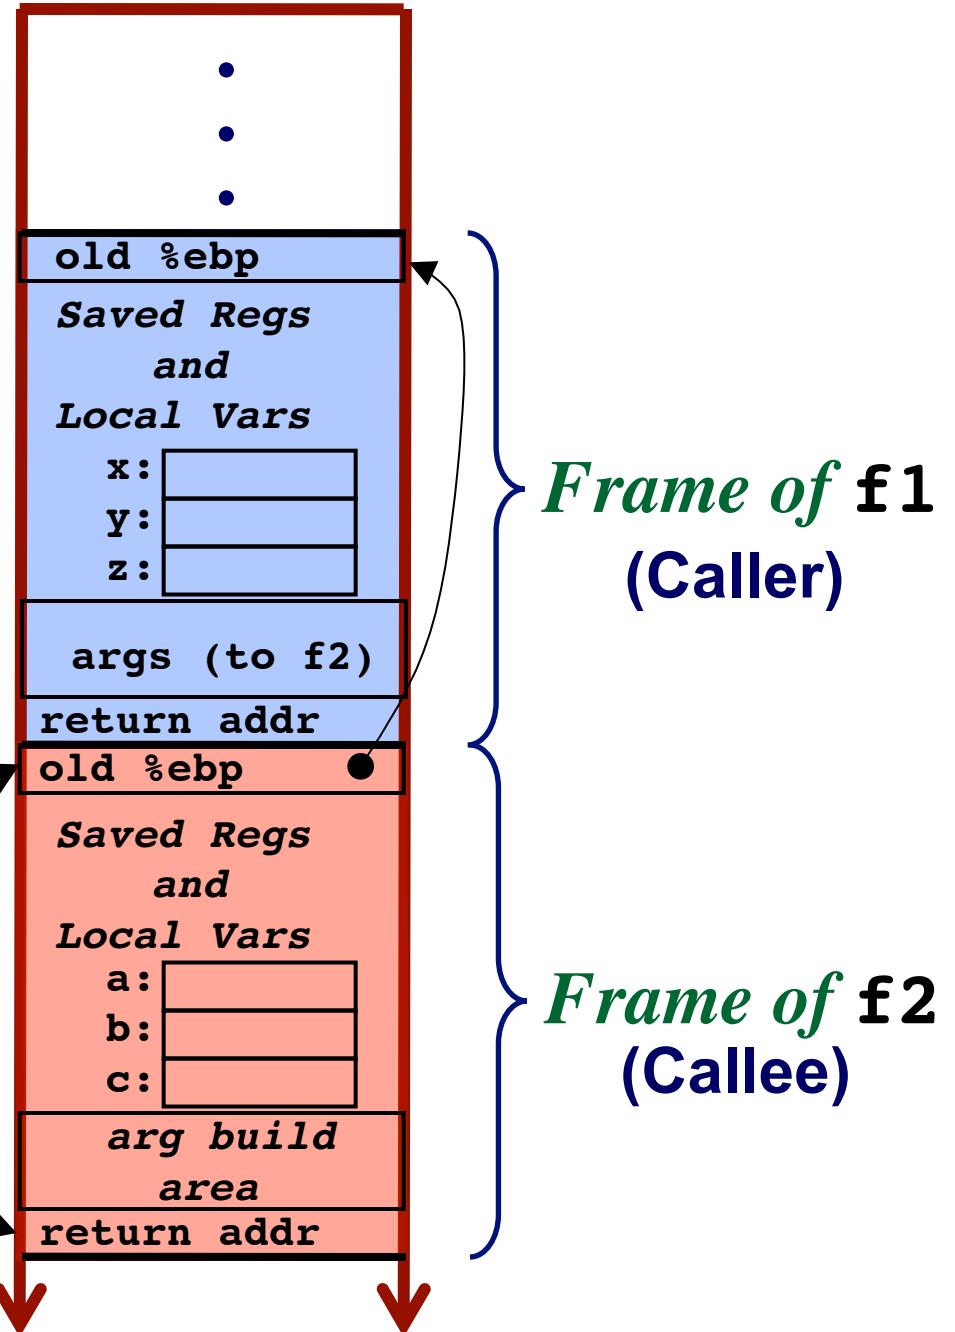
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

%esp



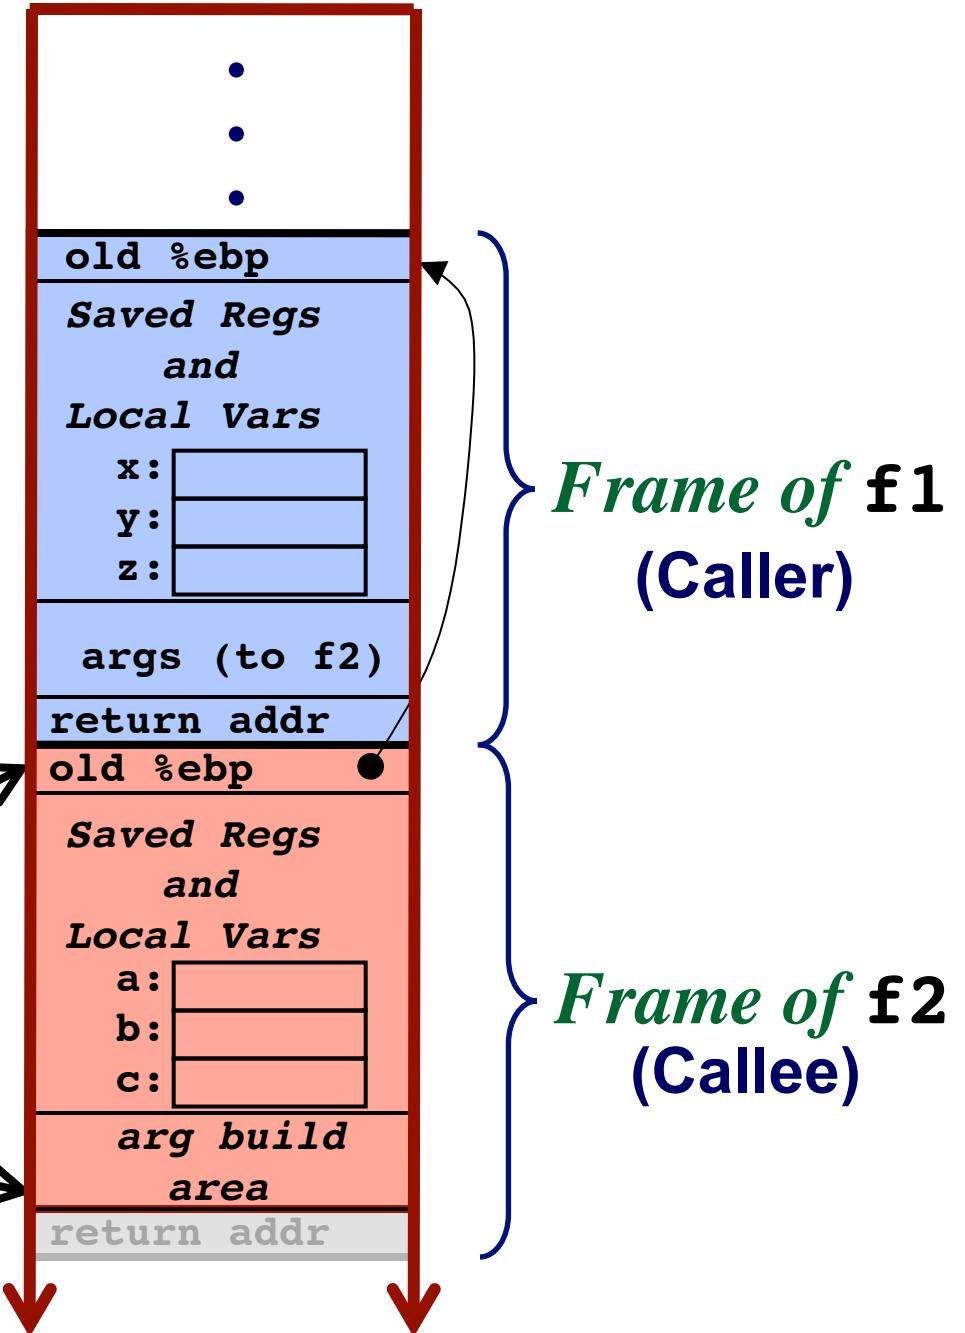
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp

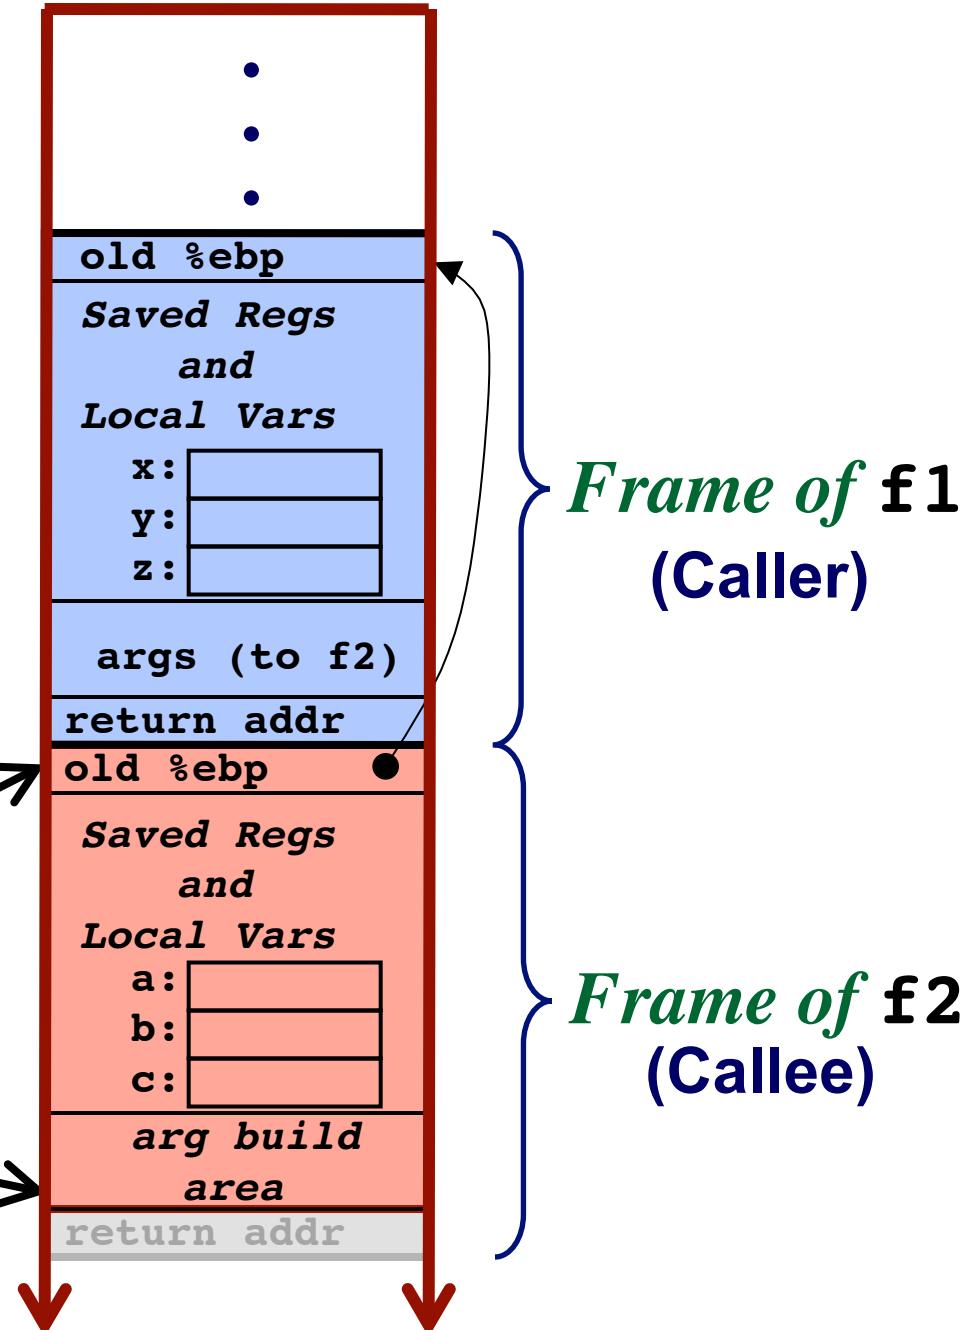
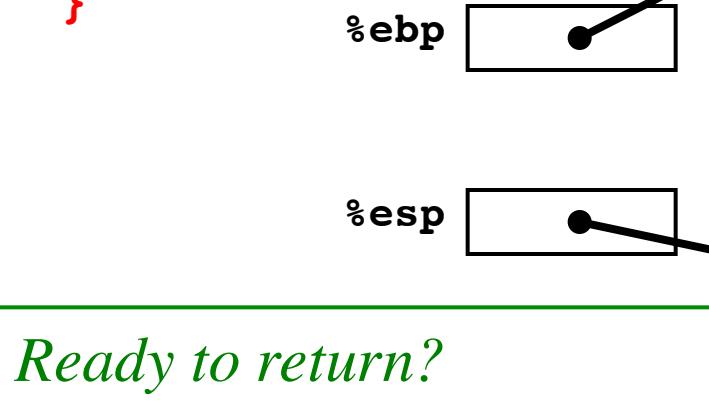
%esp



Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```



Stack Frame Layout

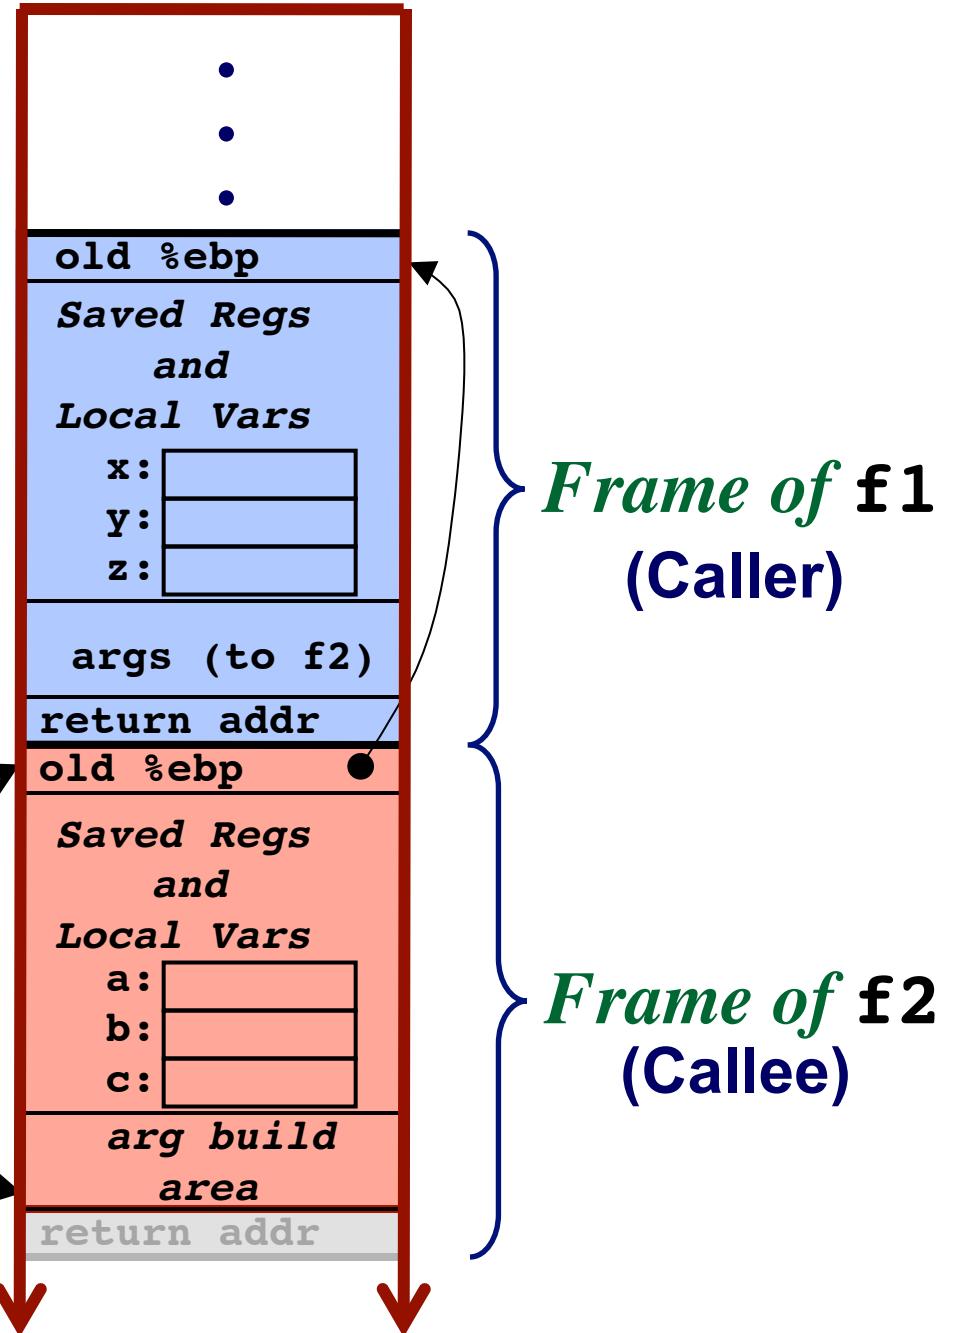
```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

`%ebp` 

`%esp` 

next instruction:
`movl %ebp, %esp`



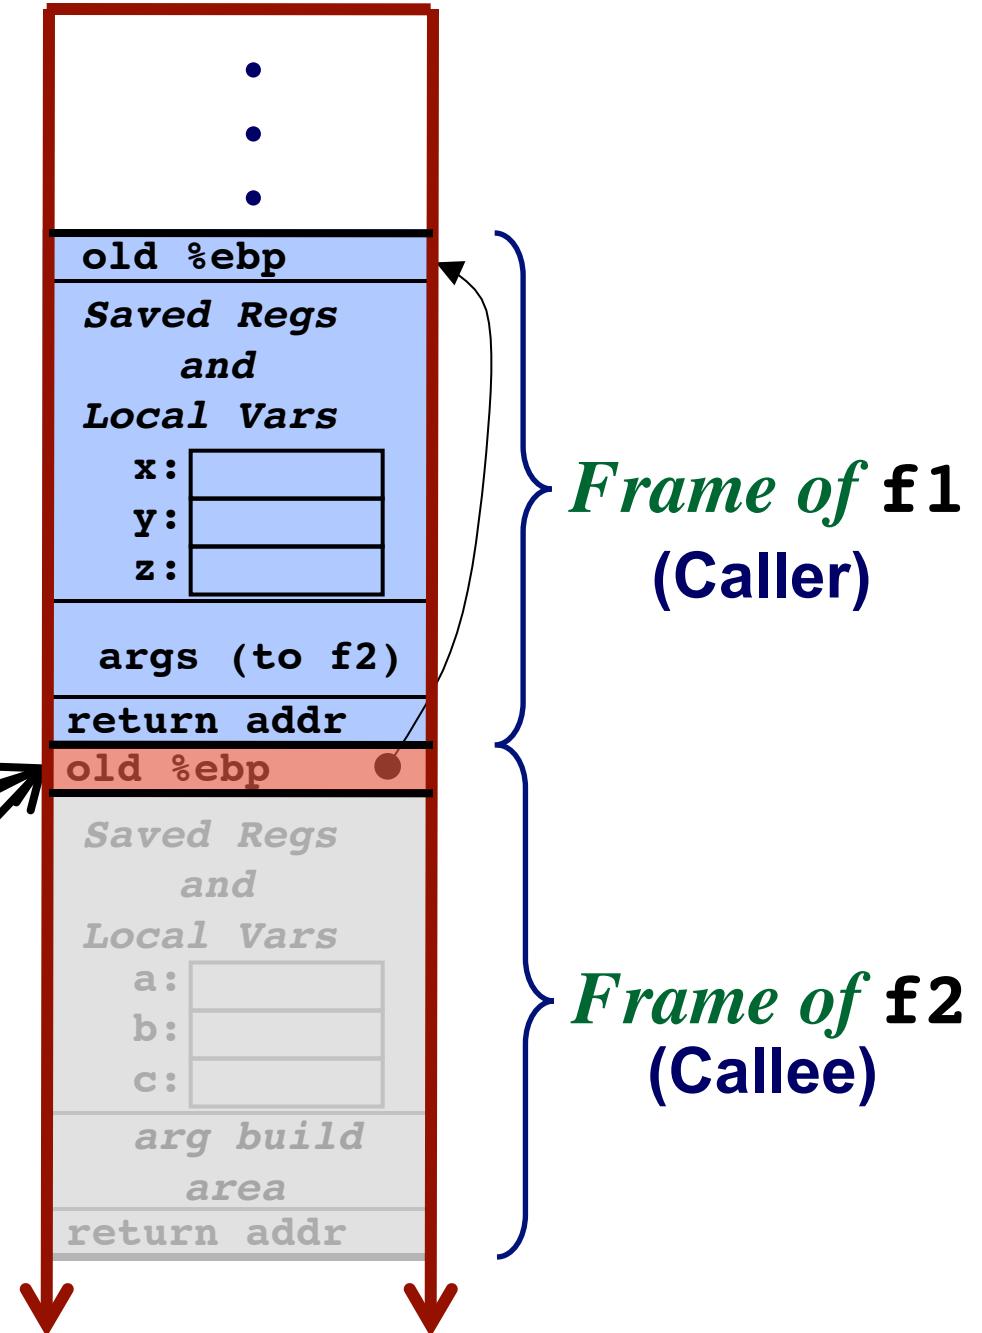
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp
%esp

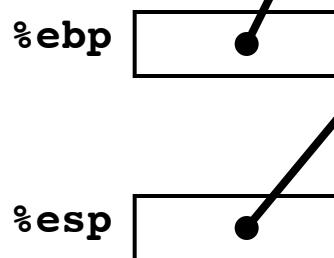
next instruction:
popl %ebp



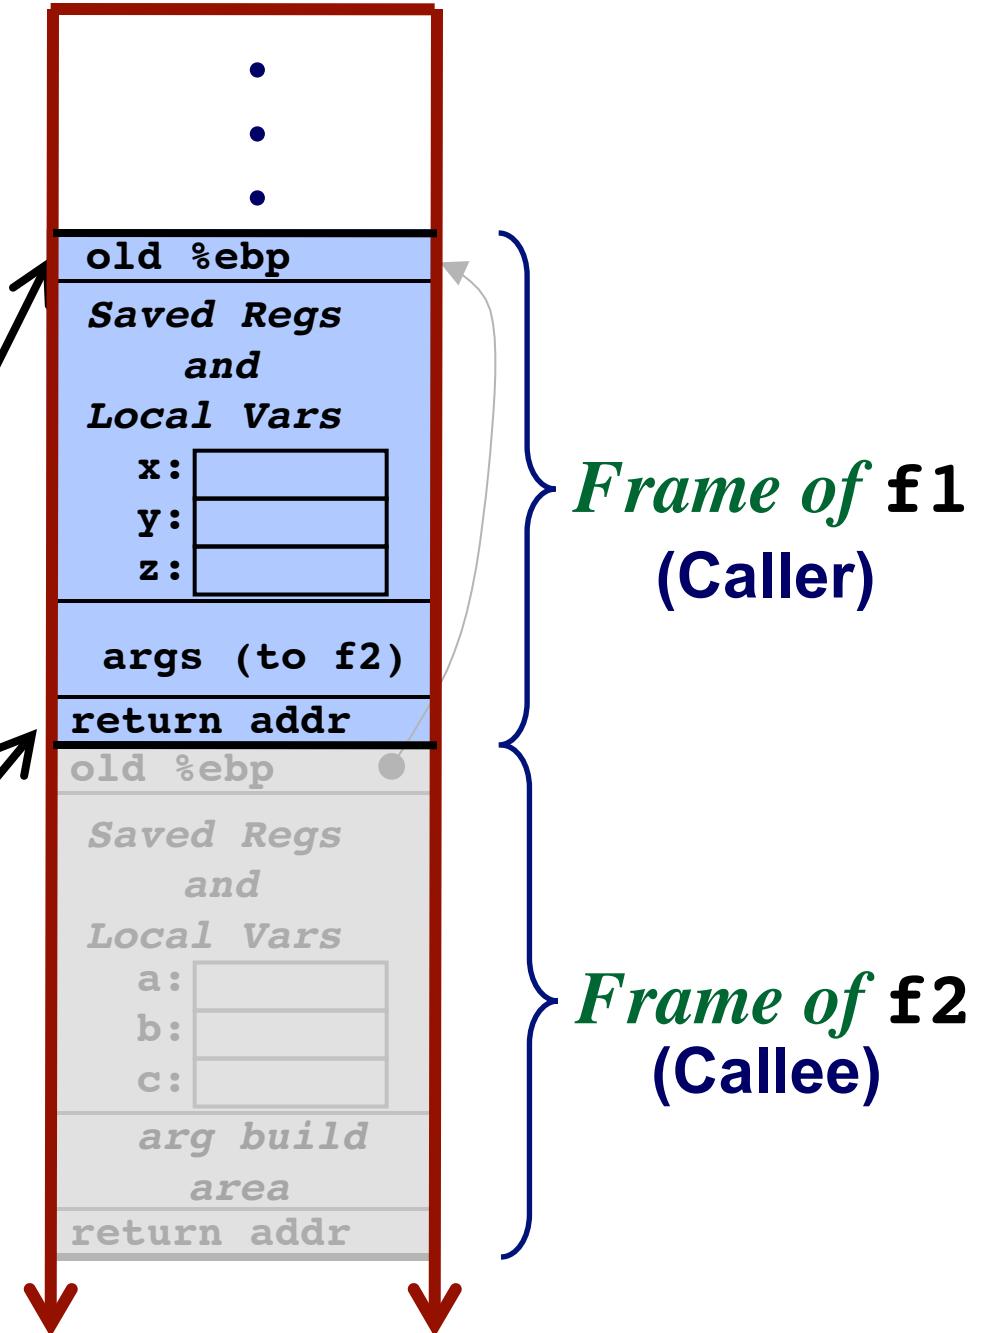
Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```



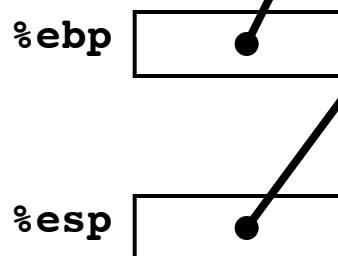
next instruction:
ret

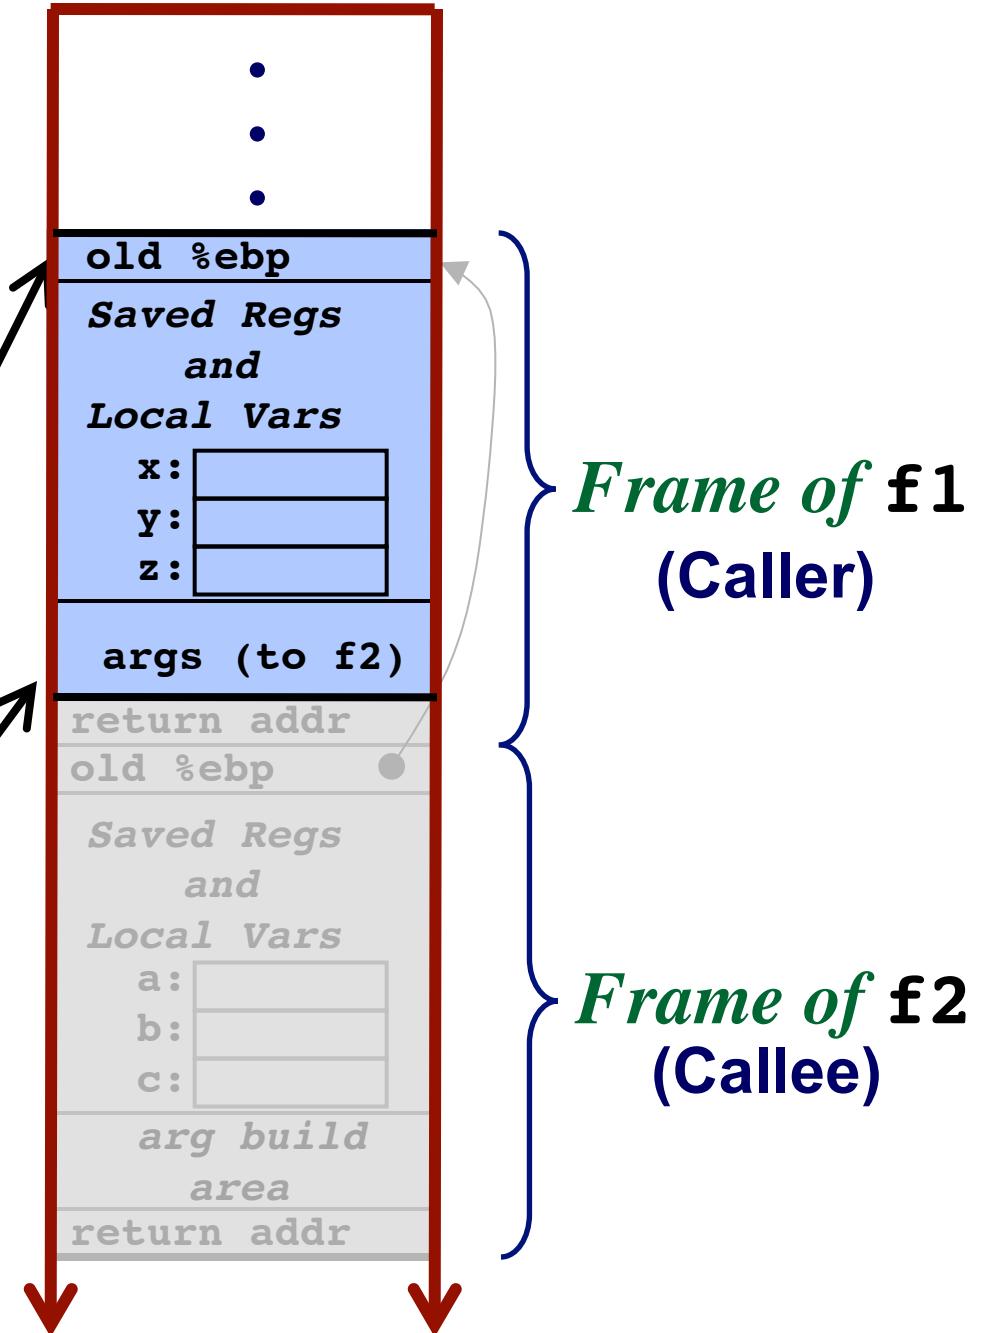


Stack Frame Layout

```
f1() {  
    int z,y,z;  
    ... call f2()  
}
```

```
f2() {  
    int a,b,c;  
    ... call f3()  
}
```

%ebp 



Revisiting swap

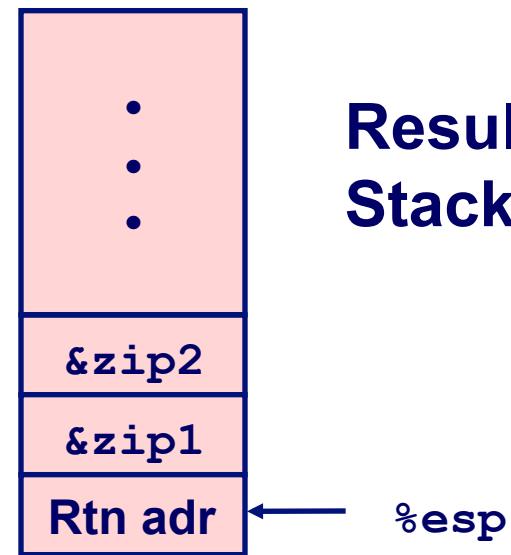
```
int zip1 = 15213;
int zip2 = 91125;

void call_swap()
{
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Calling swap from call_swap

```
call_swap:
    • • •
    pushl $zip2    # Global Var
    pushl $zip1    # Global Var
    call swap
    • • •
```



Revisiting swap

swap:

```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

} Setup

```
movl 12(%ebp),%ecx  
movl 8(%ebp),%edx  
movl (%ecx),%eax  
movl (%edx),%ebx  
movl %eax,(%edx)  
movl %ebx,(%ecx)
```

} Body

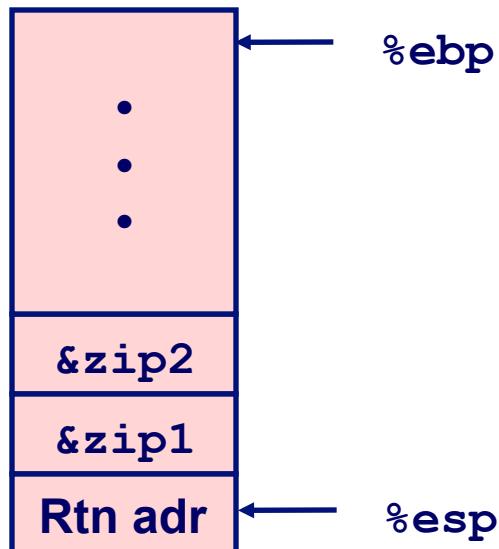
```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

} Finish

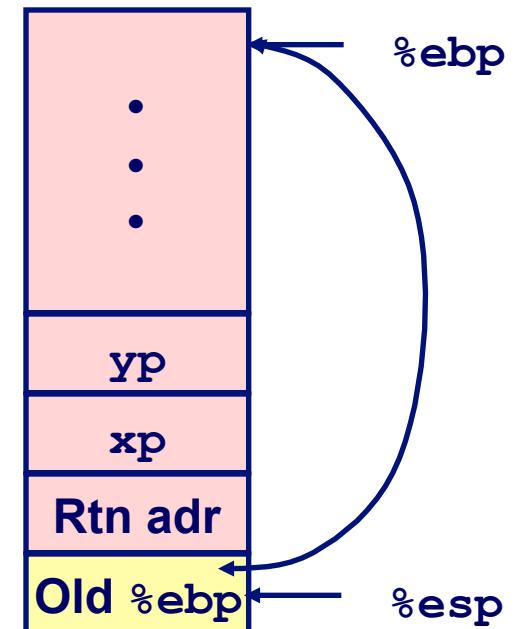
```
void swap(int *xp, int *yp)  
{  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

swap Setup #1

Entering
Stack



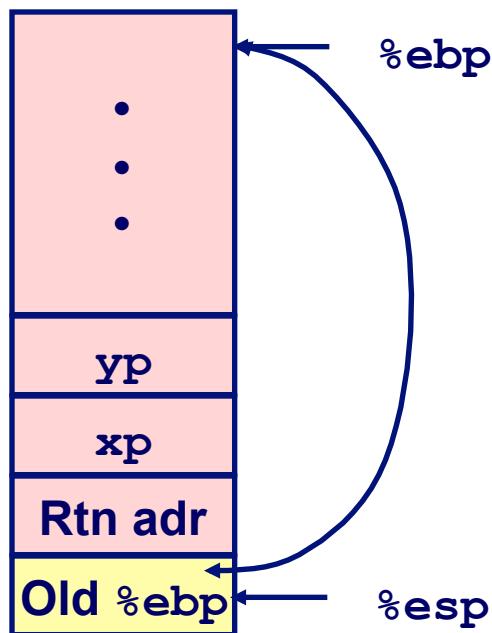
Resulting
stack



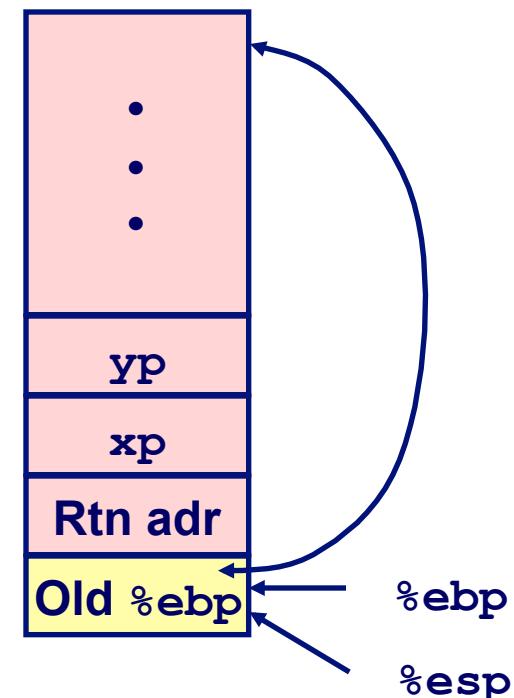
```
swap:  
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

swap Setup #2

Stack before instruction



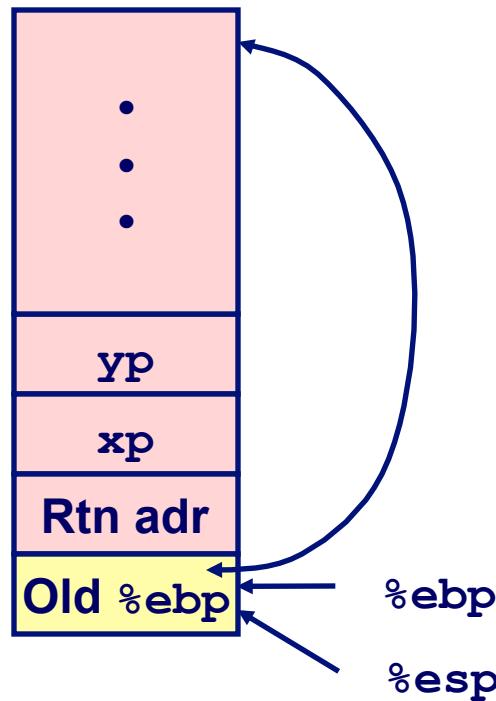
Resulting stack



```
swap:  
  pushl %ebp  
  movl %esp,%ebp  
  pushl %ebx
```

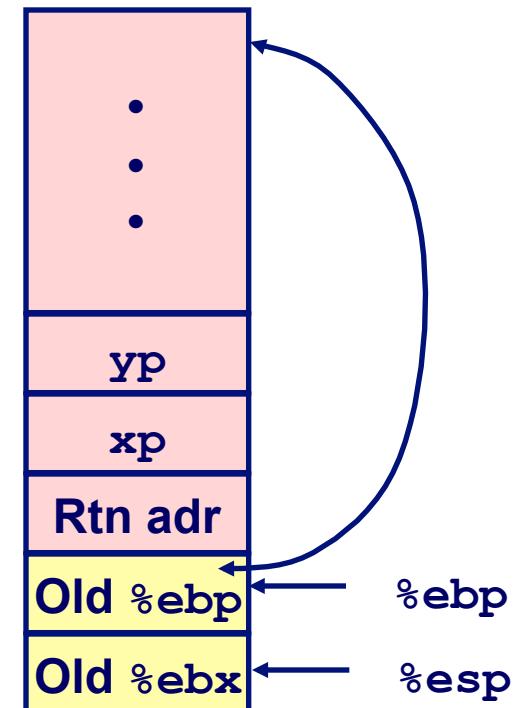
swap Setup #3

Stack before instruction



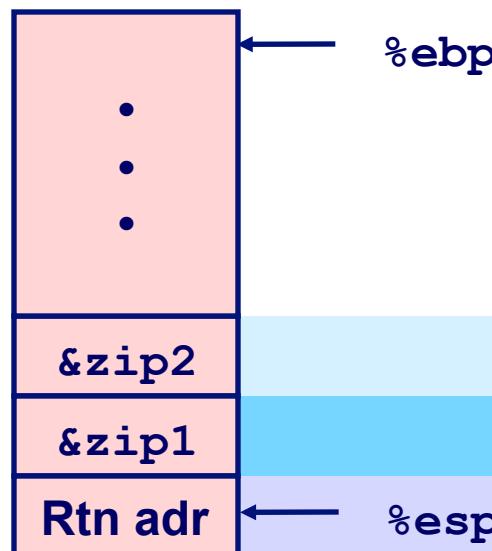
```
swap:  
  pushl %ebp  
  movl %esp,%ebp  
  pushl %ebx
```

Resulting Stack

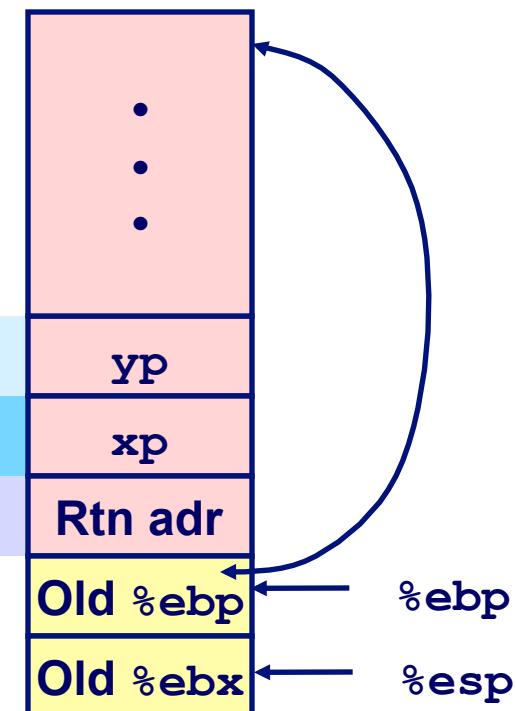


Effect of swap Setup

Entering
Stack

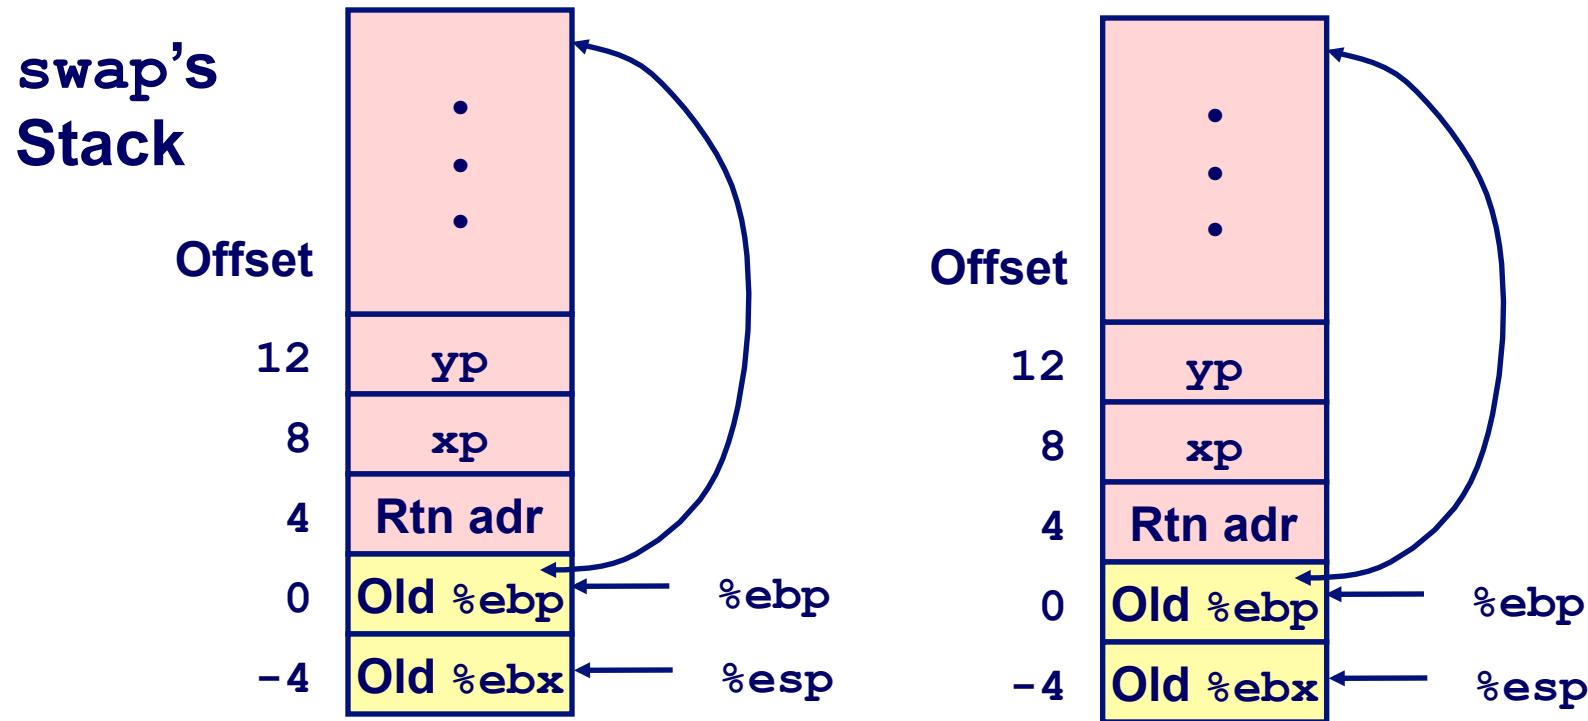


Resulting
Stack



`movl 12(%ebp),%ecx # get yp`
`movl 8(%ebp),%edx # get xp` } Body
. . .

swap Finish #1

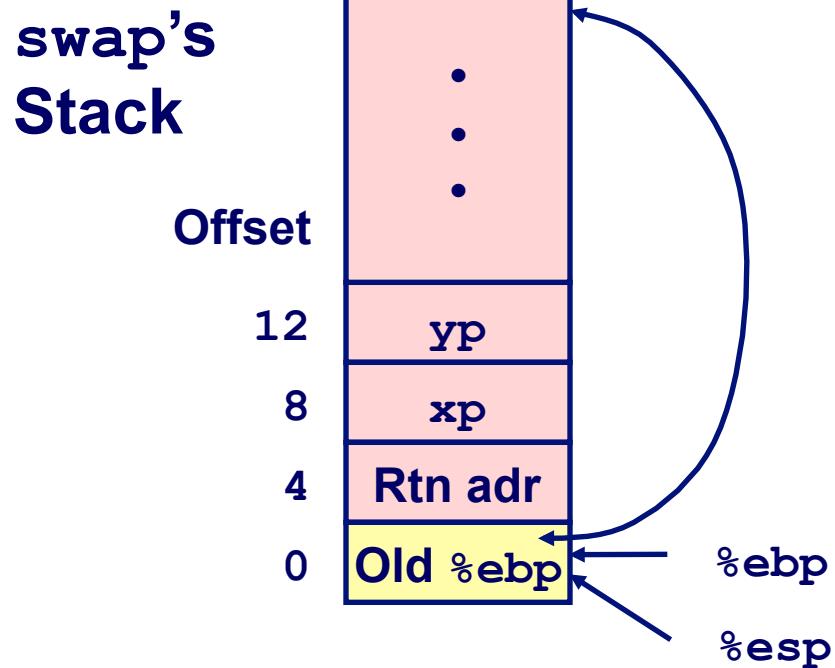
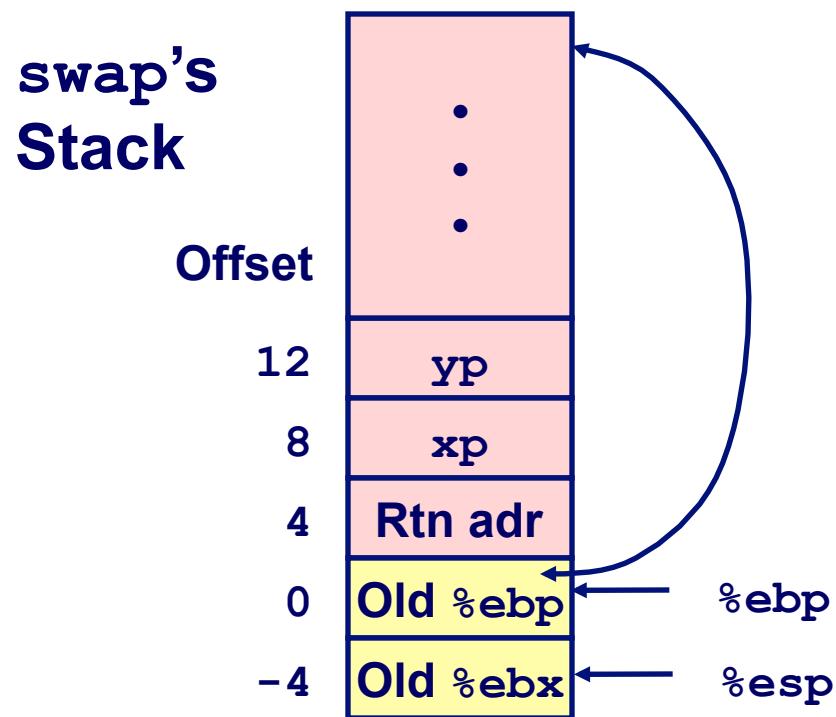


Observation

- Saved & restored register `%ebx`

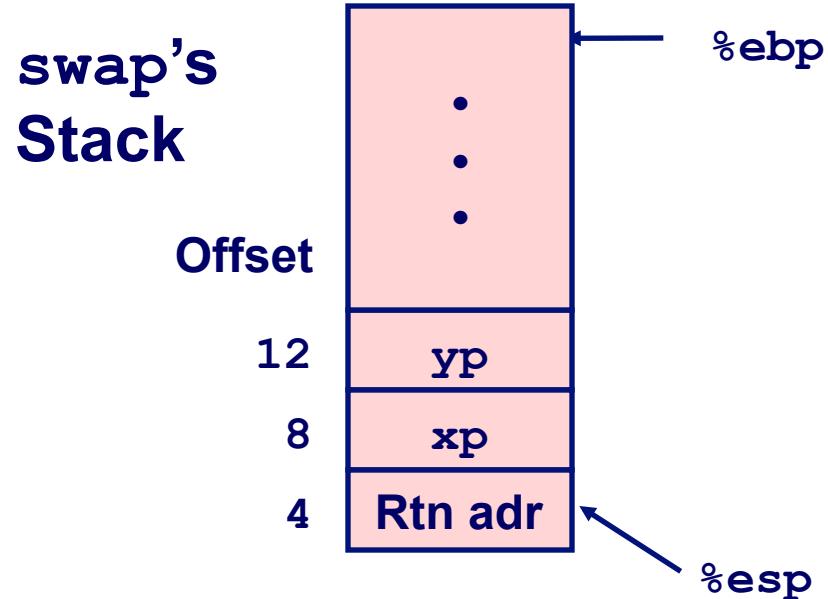
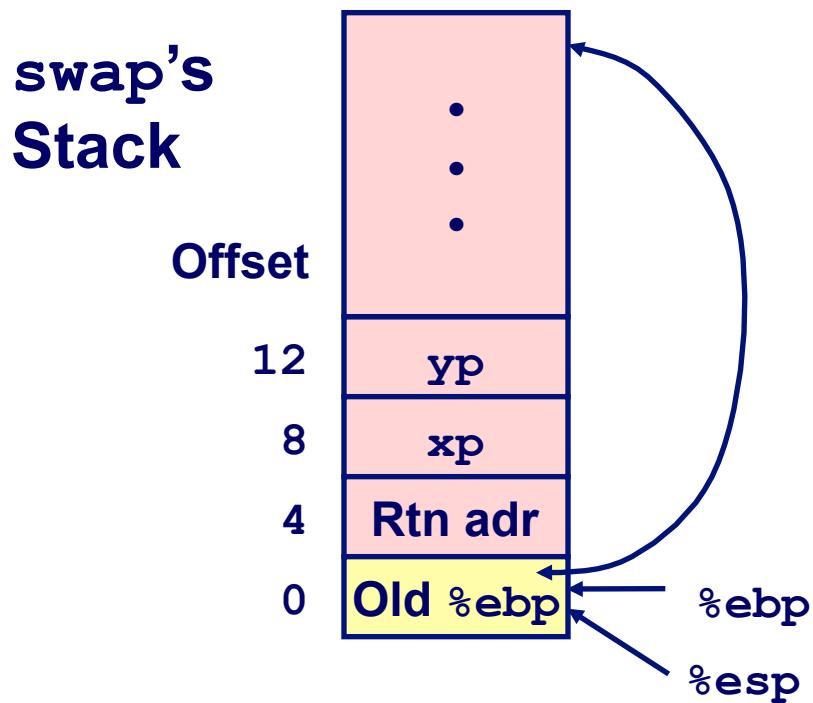
```
movl -4(%ebp),%ebx
movl %ebp,%esp
popl %ebp
ret
```

swap Finish #2



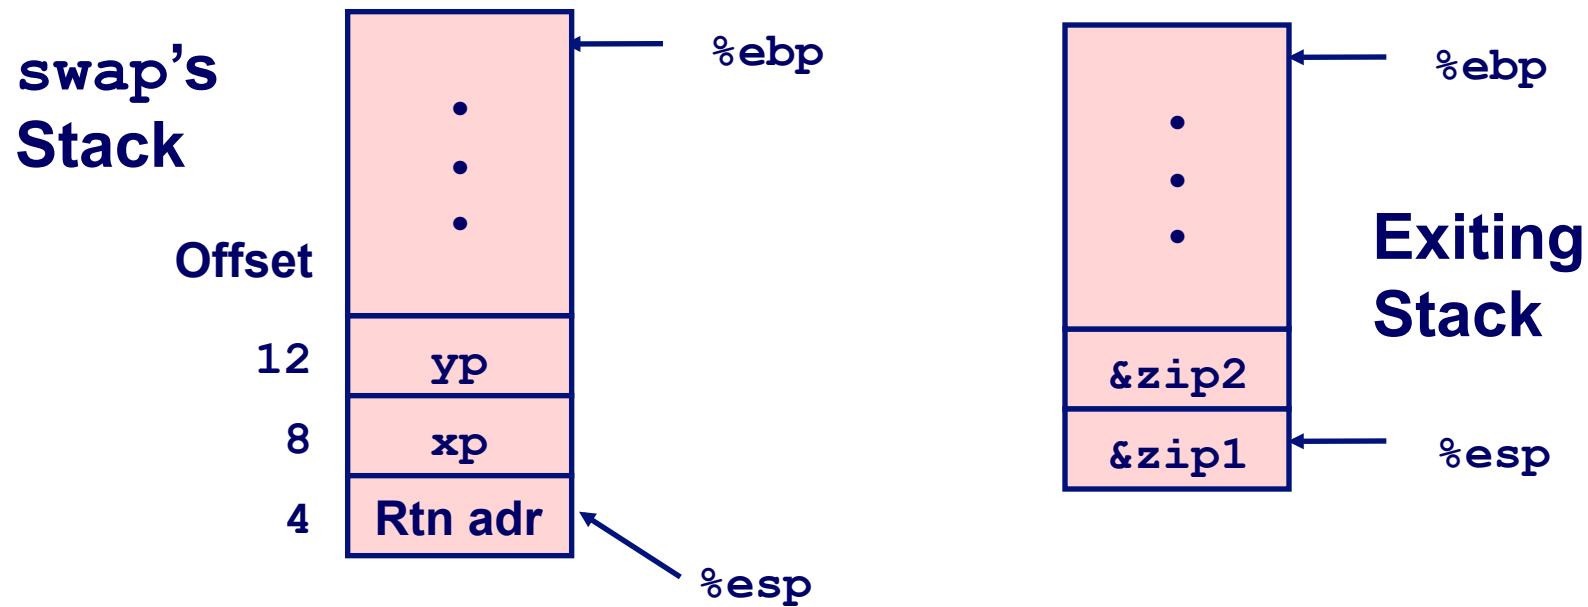
```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```

swap Finish #3



```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```

swap Finish #4



Observation

- Saved & restored register **%ebx**
- Didn't do so for **%eax**, **%ecx**, or **%edx**

```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

Revisiting swap

swap:

```
pushl %ebp
movl %esp,%ebp
pushl %ebx

movl 12(%ebp),%ecx
movl 8(%ebp),%edx
movl (%ecx),%eax
movl (%edx),%ebx
movl %eax,(%edx)
movl %ebx,(%ecx)
```

```
movl -4(%ebp),%ebx
movl %ebp,%esp
popl %ebp
ret
```

Setup

Save old %ebp of caller frame
Set new %ebp for callee (current) frame
Save state of %ebx register from caller

Body

Retrieve parameter yp from caller frame
Retrieve parameter xp from caller frame

Perform swap

Finish

Restore the state of caller's %ebx register
Set stack pointer to bottom of callee frame (%ebp)
Restore %ebp to original state
Pop return address from stack to %eip

Replaced by
leave

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Local variables

Where are they in relation to ebp?

- Stored “above” %ebp (at lower addresses)

How are they preserved if the current function calls another function?

- Compiler updates %esp beyond local variables before issuing “call”

What happens to them when the current function returns?

- Are lost (i.e. no longer valid)

What's the matter with this?

List two problems with the following code

```
int *foo(...)  
{  
    int *val;  
    ...  
    *val = 34;  
    return val;  
}
```

Exercise

```
int *func(int x)
{
    int n;
    n = x;
    return &n;
}
```

Will the C compiler give you an error on this function?

What will happen when it returns?

What if the pointer it returns is dereferenced?

Exercise

```
int *func(int x)
{
    int n;
    n = x;
    return &n;
}
```

Will the C compiler give you an error on this function?

A warning is given

What will happen when it returns?

Returns an address that is no longer part of the stack

What if the pointer it returns is dereferenced?

If pointer is outside of valid stack page, segmentation fault

Otherwise, returns whatever is at location

Register Saving Conventions

When procedure `foo` calls `who`:

`foo` is the *caller*, `who` is the *callee*

Can a Register be Used for Temporary Storage?

```
foo:  
  • • •  
  movl $15213, %edx  
  call who  
  addl %edx, %eax  
  • • •  
  ret
```

```
who:  
  • • •  
  movl 8(%ebp), %edx  
  addl $91125, %edx  
  • • •  
  ret
```

Contents of register `%edx` overwritten by `who`

Register Saving Conventions

When procedure `foo` calls who:

`foo` is the *caller*, who is the *callee*

Can a Register be Used for Temporary Storage?

Conventions

“Caller Save”

- Caller saves temporary in its frame before calling

“Callee Save”

- Callee saves temporary in its frame before using

IA32 Register Usage

Integer Registers

Two have special uses

`%ebp`, `%esp`

Three managed as callee-save

`%ebx`, `%esi`, `%edi`

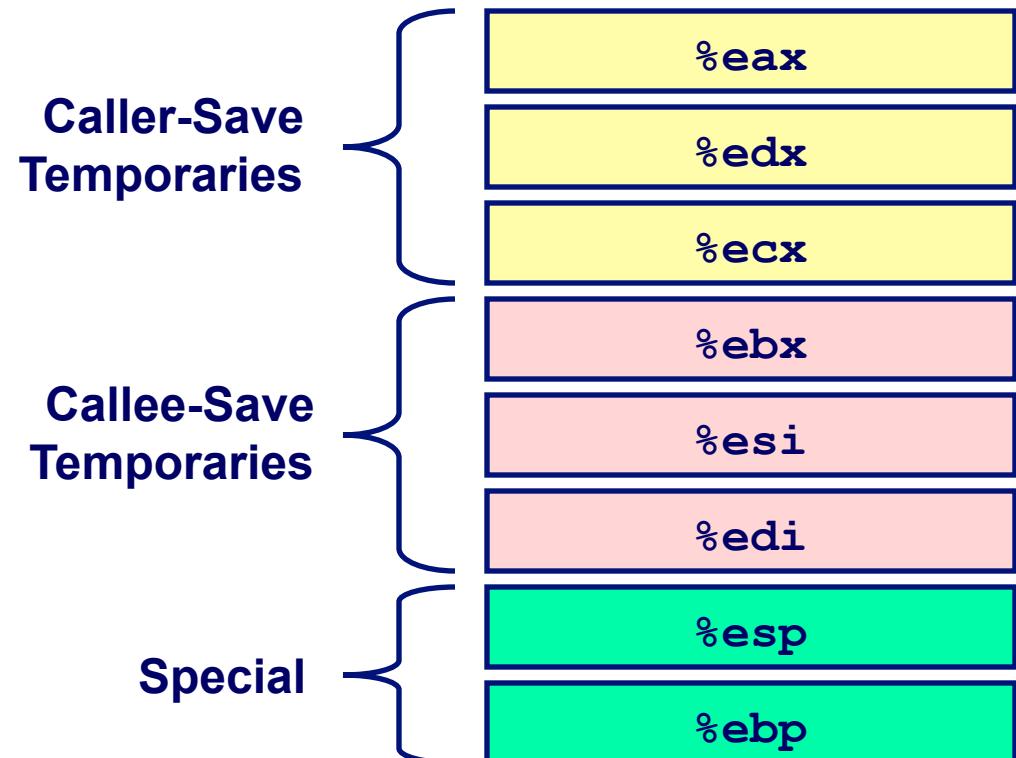
Old values should be saved on stack prior to using

Three managed as caller-save

`%eax`, `%edx`, `%ecx`

Do what you please, but expect any callee to do so, as well

Return value in `%eax`



IA32 register usage

Intel publishes these register usage conventions.

Compiler writers are expected to comply.

Why shouldn't each compiler writer make up their own conventions?

Practice problem 3.30

What does this code do?

```
    call next
next:
    popl %eax
```

What does the call do to the stack?

What is the value of %eax?

Where'd the ret go?

What would this be useful for?

Recursive Procedures

Since each call results in a new stack frame, recursive calls become natural

A recursive call is just like any other call, as far as IA32 assembly code is concerned

- Of course, the a recursive algorithm needs a termination condition, but that's the programmer's problem

<http://thefengs.com/wuchang/courses/cs201/class/08/stack.c>

Recursive Factorial

```
int rfact(int x)
{
    int rval;
    if (x <= 1)
        return 1;
    rval = rfact(x-1);
    return rval * x;
}
```

$$x! = (x-1)! * x$$

Registers

- `%eax` used without first saving
- `%ebx` saved at beginning & restored at end

`%eax` is pushed before call to `rfact` and never popped. Is that a problem?

```
.globl rfact
.type rfact,@function
rfact:
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
    movl 8(%ebp),%ebx
    cmpl $1,%ebx
    jle .L78
    leal -1(%ebx),%eax
    pushl %eax
    call rfact
    imull %ebx,%eax
    jmp .L79
    .align 4
.L78:
    movl $1,%eax
.L79:
    movl -4(%ebp),%ebx
    movl %ebp,%esp
    popl %ebp
    ret
```

Practice problem 3.33

```
int proc(void) {  
    int x,y;  
    scanf("%x %x", &y, &x);  
    return (x-y)  
}
```

```
1 proc:  
2     pushl %ebp  
3     movl %esp,%ebp  
4     subl $24,%esp  
5     addl $-4,%esp  
6     leal -4(%ebp),%eax  
7     pushl %eax  
8     leal -8(%ebp),%eax  
9     pushl %eax  
10    pushl $.LC0      /* "%x %x" */  
11    call scanf  
12    movl -8(%ebp),%eax  
13    movl -4(%ebp),%edx  
14    subl %eax,%edx  
15    movl %edx,%eax  
16    movl %ebp,%esp  
17    popl %ebp  
18    ret
```

%esp initially 0x800040
%ebp initially 0x800060
scanf reads 0x46 and 0x53 from stdin
“%x %x” stored at 0x300070

What value does %ebp get set to in line 3?

At what addresses are local variables x and y stored?

What is the value of %esp after line 10?

Practice problem 3.33

```
int proc(void) {
    int x,y;
    scanf("%x %x", &y, &x);
    return (x-y)
}
```

```
1 proc:
2     pushl %ebp
3     movl %esp,%ebp
4     subl $24,%esp
5     addl $-4,%esp
6     leal -4(%ebp),%eax
7     pushl %eax
8     leal -8(%ebp),%eax
9     pushl %eax
10    pushl $.LC0      /* "%x %x" */
11    call scanf
12    movl -8(%ebp),%eax
13    movl -4(%ebp),%edx
14    subl %eax,%edx
15    movl %edx,%eax
16    movl %ebp,%esp
17    popl %ebp
18    ret
```

%esp initially 0x800040
 %ebp initially 0x800060
 scanf reads 0x46 and 0x53 from stdin
 "%x %x" stored at 0x300070

What value does %ebp get set to in line 3?

0x80003C

At what addresses are local variables x and y stored?

0x800038 and 0x800034

What is the value of %esp after line 10?

$0x800040 - 4 - 24 - 4 - 4 - 4 - 4$
 $0x800040 - 0x2C = 0x800014$

Practice problem 3.33

```
int proc(void) {
    int x,y;
    scanf("%x %x", &y, &x);
    return (x-y)
}
```

```
1 proc:
2     pushl %ebp
3     movl %esp,%ebp
4     subl $24,%esp
5     addl $-4,%esp
6     leal -4(%ebp),%eax
7     pushl %eax
8     leal -8(%ebp),%eax
9     pushl %eax
10    pushl $.LC0      /* "%x %x" */
11    call scanf
12    movl -8(%ebp),%eax
13    movl -4(%ebp),%edx
14    subl %eax,%edx
15    movl %edx,%eax
16    movl %ebp,%esp
17    popl %ebp
18    ret
```

%esp initially 0x800040
 %ebp initially 0x800060
 scanf reads 0x46 and 0x53 from stdin
 "%x %x" stored at 0x300070

Draw a diagram of the stack frame for proc right after scanf returns and indicate the regions of the stack frame that are not used.

0x80003C	0x800060
0x800038	0x53
0x800034	0x46
0x800030	
0x80002C	
0x800028	
0x800024	
0x800020	
0x80001C	0x800038
0x800018	0x800034
0x800014	0x300070
0x800010	

Other calling conventions

Previous slides describe the “cdecl” calling convention

- Two other calling conventions “stdcall” and “fastcall”

cdecl

- Caller pushes arguments on stack before call
- Caller clears arguments off stack after call

stdcall

- Standard Win 32 API
- Same as cdecl except ...
- Caller pushes arguments on stack before call
- Callee clears arguments off stack before returning from call
 - Saves some instructions since callee is already restoring the stack at the end of the function

Other calling conventions

fastcall

- Save memory operations by passing arguments in registers
- Many versions based on the compiler
 - Combining code that has been compiled with different conventions ugly
 - Must declare in source code if calling convention is different for function
- Microsoft implementation
 - First two arguments passed in registers %ecx and %edx
 - Code written on Windows must deal with stdcall and fastcall conventions
- Linux
 - Must declare in function prototype which calling convention is being used
 - <http://gcc.gnu.org/onlinedocs/gcc/Function-Attributes.html>

x86-64

- Function arguments placed in registers
 - 4 integer register arguments
 - 4 floating point register arguments
 - Rest on stack

Other calling conventions

thiscall

- Used for C++
- Linux
 - Same as cdecl, but first argument assumed to be “this” pointer
- Windows/Visual C++
 - “this” pointer passed in %ecx
 - Callee cleans the stack when arguments are not variable length
 - Caller cleans the stack when arguments are variable length

More information

- <http://www.programmersheaven.com/2/Calling-conventions>

Extra slides

Summary

The Stack Makes Recursion Work

- Private storage for each *instance* of procedure call
 - Instantiations don't clobber each other
 - Addressing of locals + arguments can be relative to stack positions
- Can be managed by stack discipline
 - Procedures return in inverse order of calls

IA32 Procedures Combination of Instructions + Conventions

- Call / Ret instructions
- Register usage conventions
 - Caller / Callee save
 - %ebp and %esp
- Stack frame organization conventions

The Towers of Hanoi

```
hanoi (int N, int src, int aux, int dst)
{
    if (N == 0) return;
    hanoi (N-1, src, dst, aux);
    move (src, dst);
    hanoi (N-1, aux, src, dst);
}
```

where “move” does the graphical or physical moving of a ring to another post.

Does this code really work?